

Contents

1	Introduction	4
1.1	Path Integral Molecular Dynamics	4
1.2	Implementation	5
1.3	Core features	8
1.4	Licence and credits	10
1.5	On-line resources	10
2	Getting started	11
2.1	Installing i-PI	11
2.2	Installing clients	13
2.3	Running i-PI	13
2.4	Testing the install	16
3	User guide	17
3.1	Input files	17
3.2	Output files	19
3.3	Distributed execution	28
4	A simple tutorial	32
4.1	Part 1 - <i>NVT</i> Equilibration run	32
4.2	Part 2 - <i>NPT</i> simulation	40
4.3	Part 3 - A fully converged simulation	41
5	Input file reference	42
5.1	simulation	42
5.2	system	44
5.3	system_template	45
5.4	motion	46
5.5	cell	47
5.6	smotion	48
5.7	remd	49
5.8	metad	49
5.9	dmd	50
5.10	ensemble	50
5.11	bias	51
5.12	dynamics	51
5.13	constrained_dynamics	52

5.14	csolver	53
5.15	constraint	54
5.16	t_ramp	55
5.17	p_ramp	56
5.18	alchemy	56
5.19	planetary	57
5.20	atomswap	59
5.21	instanton	59
5.22	vibrations	63
5.23	scp	64
5.24	normalmodes	66
5.25	optimizer	69
5.26	neb_optimizer	71
5.27	string_optimizer	74
5.28	thermostat	78
5.29	barostat	80
5.30	h0	81
5.31	forcefield	82
5.32	ffsocket	83
5.33	fflj	84
5.34	ffdebye	85
5.35	ffplumed	87
5.36	ffyaff	88
5.37	ffsgdml	90
5.38	ffdmd	91
5.39	ffcommittee	93
5.40	ffcavphsocket	95
5.41	forces	97
5.42	force	98
5.43	al6xxx_kmc	99
5.44	atoms	101
5.45	beads	102
5.46	prng	103
5.47	normal_modes	104
5.48	frequencies	105
5.49	initialize	106
5.50	file	107
5.51	positions	108
5.52	momenta	108
5.53	labels	109
5.54	masses	110
5.55	velocities	110
5.56	init_cell	111
5.57	gle	111
5.58	output	112
5.59	properties	112
5.60	checkpoint	113
5.61	trajectory	114
6	Internal units	115
7	Troubleshooting	115
7.1	Input errors	115
7.2	Initialization errors	116

7.3	Output errors	117
7.4	Socket errors	118
7.5	Mathematical errors	119
7.6	I-PI is slow!	119
8	Frequently asked questions	119
8.1	Where to ask for help?	119
8.2	Which units does i-PI use?	120
8.3	How to build i-PI?	120
8.4	How to run i-PI with the client code <name>?	120
8.5	How to run i-PI on a cluster?	121
8.6	How to setup colored-noise thermostats with meaningful parameters?	121
8.7	How to perform geometry optimization?	121
9	Contributing	121
9.1	Useful resources	121
9.2	What we expect from a contribution	121
9.3	Creating a regression test	122
9.4	Adding features involving a client code	122
9.5	Getting recognition for your contribution	123
9.6	Contributing with a bug report	123
9.7	Questions related to i-PI usage	123
10	Bibliography	123
	References	123

i-PI is a interface for advanced molecular simulations written in Python, designed to be used together with an *ab initio* evaluation of the interactions between the atoms. The main goal is to decouple the problem of evolving the ionic positions to sample the appropriate thermodynamic ensemble and the problem of computing the inter-atomic forces. i-PI was initially developed for Path Integral Molecular Dynamics (PIMD) simulations, and contains probably the most comprehensive array of PIMD techniques. Since v2.0, however, it also contains several general-purpose methods for molecular simulations, ranging from phonon calculators to replica exchange molecular dynamics.

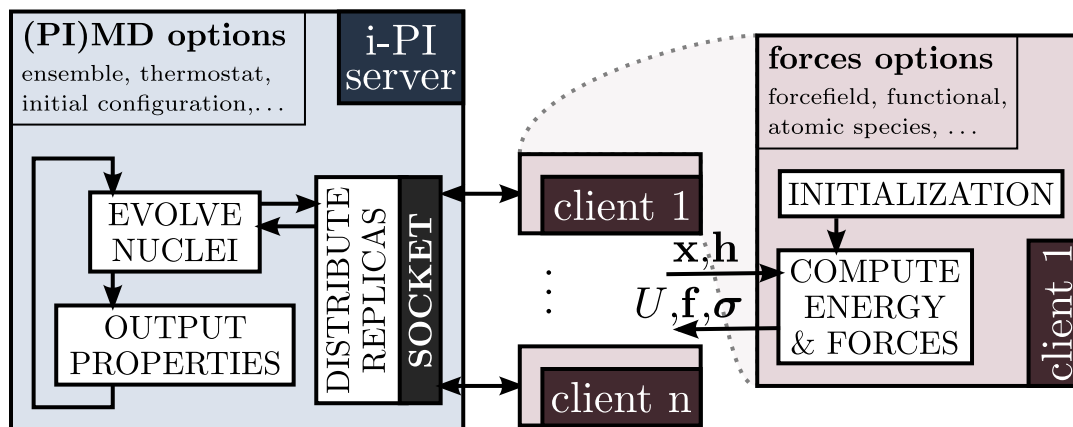


Fig. 1: Schematic representation of the functioning of i-PI.

The implementation is based on a client-server paradigm, where i-PI acts as the server and deals with the propagation

of the nuclear dynamics, whereas the calculation of the potential energy, forces and the potential energy part of the pressure virial is delegated to one or more instances of an external code, acting as clients. Since the main focus is on performing *ab initio* PIMD – where the cost of the force evaluation is overwhelming relative to the ionic dynamics – clarity has been privileged over speed. Still, the implementation of i-PI is efficient enough that, by tuning the socket parameters and avoiding excessive I/O it can be used to run empirical forcefields with tens of thousands of atoms with only a small overhead.

This documentation is structured as follows:

1 Introduction

1.1 Path Integral Molecular Dynamics

Molecular dynamics (MD) is a technique used to study the properties of a system of interacting particles by applying Newton’s equations of motion to produce trajectories which can be used to efficiently explore the phase space. This can be used to calculate many equilibrium and dynamical properties and to study systems from isolated gas molecules to condensed phase bulk materials.

However, while this technique has been very successful, in most MD implementations the assumption is made that the nuclei behave as classical particles, which for light nuclei such as hydrogen is often a very poor approximation as the effect of zero-point energy (ZPE) and quantum tunnelling can be large. For example, even at room temperature the vibrational frequency of an OH stretch in water is over 15 times larger than the available thermal energy, and so this motion will be highly quantized. The current state-of-the-art method to include nuclear quantum effects (NQE) in the calculation of static properties of condensed phase systems is path integral molecular dynamics (PIMD).

PIMD generates the quantum-mechanical ensemble of a system of interacting particles by using MD in an extended phase space. This is derived from the path integral formalism [FH64], which relates the statistics of a collection of quantum particles to those of a set of classical ring polymers, a ring polymer being a number of replicas of a particle coupled by harmonic springs. This so-called classical isomorphism is exact in the limit as the number of replicas goes to infinity, but in practice is converged numerically with only a finite number.

This then allows quantum phase space averages to be calculated from classical trajectories, with only about an order of magnitude more computing time than would be required for standard MD. Also, since PIMD is simply classical MD in an extended phase space, many of the techniques developed to improve the scope and efficiency of MD simulations can be applied straightforwardly to the equivalent PIMD calculations [CPMM10][MHT99]. Finally, several techniques designed specifically for PIMD simulations are now available to increase the rate of convergence with respect to the number of replicas used [MM08b][CMP11][Suz95][Chi97][CBRM11][PerezT11], further reducing the computational overhead of the method. All of these facts mean that it is now feasible to do PIMD simulations with thousands of molecules, or even to use *ab initio* electronic structure calculations to propagate the dynamics for small systems.

Furthermore, the framework used to run PIMD simulations can be adapted to generate approximate quantum dynamical information [CV93][CV94][CM04][BM06], and so can also be used to calculate correlation functions. While real-time quantum coherences cannot be captured, the inclusion of quantum statistical information and the rapid decoherence observed in condensed phase systems mean that in many cases very accurate results can be obtained from such approximate treatments of quantum dynamics [HMMM13].

1.2 Implementation

Automated evaluation (depend objects)

i-PI uses a caching mechanism with automatic value updating to make the code used to propagate the dynamics as simple and clear as possible. Every physical quantity that is referenced in the code is created using a “depend” object class, which is given the parameters on which it depends and a function used to calculate its value.




Fig. 2: Schematic overview of the functioning of the *depend* class used as the base for properties and physical quantities in i-PI. A few “primitive” quantities – such as atomic positions or momenta – can be modified directly. For most properties, one defines a function that can compute that property based on the value of other properties. Whenever one property is modified, all the quantities that depend on it are marked as tainted, so that when the value of one of the properties is used, the function can be invoked and the updated value obtained. If a quantity is not marked as tainted, the cached value is returned instead.

“Depend” objects can be called to get the physical quantity they represent. However, they have further functionality. Firstly, once the value of a “depend” object has been calculated, its value is cached, so further references to that quantity will not need to evaluate the function that calculates it. Furthermore, the code keeps track of when any of the dependencies of the variable are updated, and makes sure that the quantity is automatically when it is needed (i.e., when the quantity is assessed again).

This is a minimal example of how to implement dependencies in a class

```
from ipi.utils.depend import depend_value, dproperties

class DObject:
    def __init__(self):
        # depend objects are created using an underscore prefix.
        # this is a "primitive" value that doesn't depend on anything
        self._scalar = depend_value(value=1, name="scalar")

        # this is a dependent object. the definition contains a function that
        # is called to determine the value, and specification of what objects
        # it depend on.
        self._double = depend_value(func=lambda: 2*self._scalar, name="double",
                                    dependencies=[self._scalar])

    # after the definition of a class, this helper function should be called to
    # create property getters and setters (use the names with no leading underscore)
    # note that property accessors are added to the class, not to the instances
    dproperties(DObject, ["scalar", "double"])
```

(continues on next page)

```

myobj = DObject()
# "primitive values" can be set manually
myobj.scalar = 4
# dependent objects will be computed automatically on demand
print(myobj.double) # prints `8`

```

This choice makes implementation slightly more complex when the physical observables are first introduced as variables, as one has to take care of stating their dependencies as well as the function that computes them. However, the advantage is that when the physical quantities are used, in the integrator of the dynamics or in the evaluation of physical properties, one does not need to take care of book-keeping and the code can be clean, transparent and readable.

It is also possible to define dependencies between different objects, in which case it's necessary to make sure that the compute function has access, at runtime, to the value of the other object. A typical usage pattern is

```

# NB: this is meant to be run after the previous code snippet
from ipi.utils.depend import depend_array
import numpy as np

class DOther:
    def __init__(self):

        # depend arrays must be initialized with storage space
        self._vec = depend_array(value=np.ones(4), name="vec")

    def bind(self, factor):

        self.factor = factor # stores a reference to the object holding the value
        self._scaled = depend_array(value=np.ones(4), name="vec",
                                     func=self.get_scaled, dependencies=[self._vec])

        # dependencies (or dependants) can also be added after object creation
        self._scaled.add_dependency(self.factor._double)

    def get_scaled(self):
        # computes a scaled version of the vector
        return self.vec*self.factor.double

dproperties(DOther, ["vec", "scaled"])

myoth = DOther() # creates the object
myoth.bind(myobj) # makes connections

myoth.vec = np.asarray([0,1,2,3])
print(myoth.scaled) # prints [0,8,16,24]

myoth.vec[3] = 0 # depend_arrays can be accessed as normal np.ndarray
print(myoth.scaled) # prints [0,8,16,0]

```

Force evaluation

Within i-PI, the evaluation of the forces plays a crucial role, as it is the step requiring communication with the client code. In order to have a flexible infrastructure that makes it possible to perform simulations with advanced techniques such as ring-polymer contraction [MM08b], the force evaluation machinery in i-PI might appear complicated at first, and deserves a brief discussion.

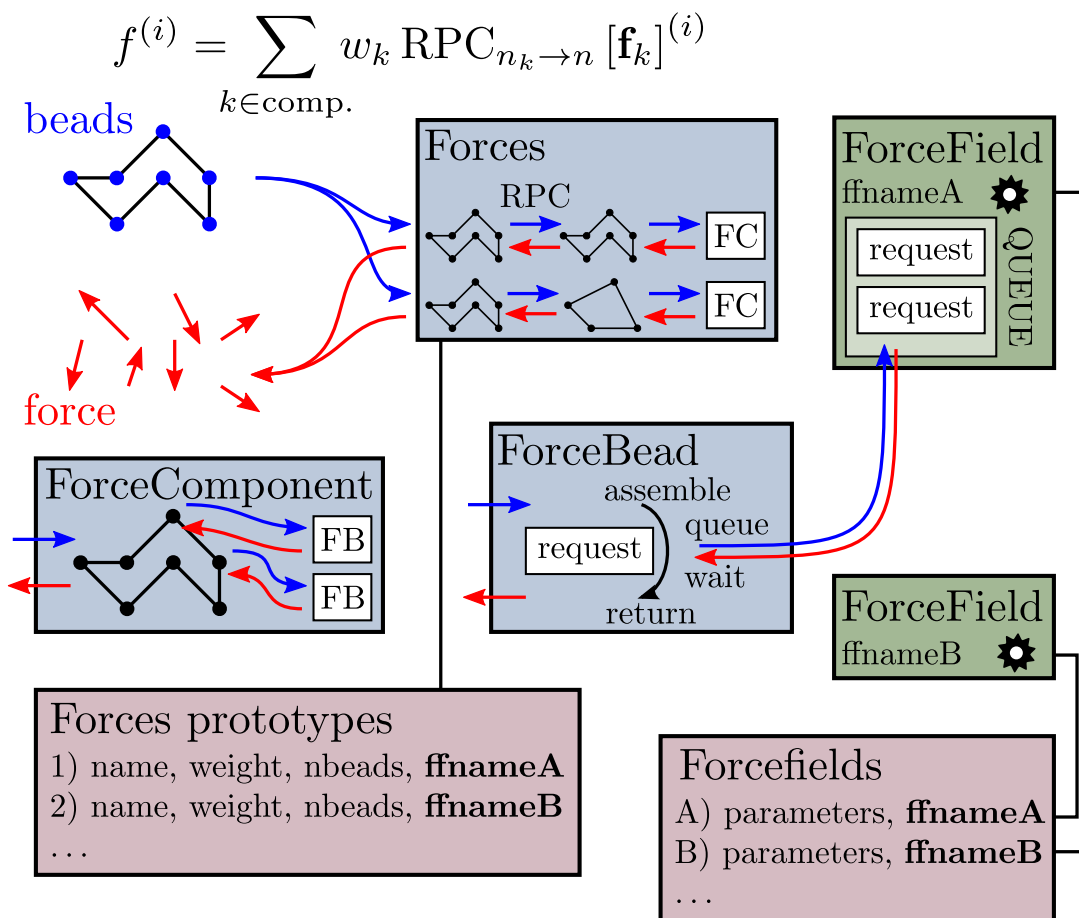


Fig. 3: Schematic representation of the different objects that are involved in the evaluation of the forces. The multiple layers and complex structure are necessary to give the possibility of decomposing the evaluation of the forces between multiple different clients and using different imaginary time partitioning (e.g. one can compute the bonded interactions using one client, and use a different client to compute the long-range electrostatic interactions, contracted on a single bead [MM08b]).

A scheme of the objects involved in the calculation of the forces is presented in Figure 1.3. The infrastructure comprises a force provider class that deals with the actual subdivision of work among the clients, and a sequence of objects that translate the request of the overall force of the system into atomic evaluations of one component of the force for an individual bead: i-PI is built to hide the path integral infrastructure from the client, and so beads must be transferred individually.

Let us discuss for clarity a practical example: a calculation of an empirical water model where the bonded interactions are computed on 32 beads by the program A, and the non-bonded interactions are computed by client B, ring-polymer contracted on 8 beads. Each client “type” is associated with a object in the input. In the case of a interface, the forcefield object specifies the address to which a client should connect, and so multiple clients of type A or B can connect to i-PI at the same time. Each forcefield object deals with queueing force evaluation requests and computing them in a

first-in-first-out fashion, possibly executing multiple requests in parallel.

On the force evaluation side, the task of splitting the request of a force evaluation into individual components and individual beads is accomplished by a chain of three objects, `Forces`, `ForceComponent` and `ForceBead`. `Forces` is the main force evaluator, that is built from the prototypes listed within the field of `ForceField`. Each item within the tag `ForceField` describes one component of the force – in our example one `ForceComponent` bound to a forcefield of type A, evaluated on 32 beads, and one `ForceComponent` bound to type B, evaluated on 8 beads. `Forces` contains the machinery that automatically contracts the actual ring polymer to the number of beads required by each component, and combines the various components with the given weights to provide the overall force, energy and virial where required. Note that in order to support ring polymer contraction (RPC), the RPC procedure is executed even if no contraction was required (i.e. even if all clients contain the full amount of beads). `ForceComponent` is a very simple helper class that associates with each bead a `ForceBead` object, that is the entity in charge of filing force requests to the appropriate `ForceField` object and waiting for completion of the evaluation.

Communication protocol

Since i-PI is designed to be used with a wide range of codes and platforms, it has to rely on a simple and robust method for communicating between the server and client. Even though other choices are possible, and it should be relatively simple to implement other means of communication, the preferred approach relies on sockets as the underlying infrastructure. Both Internet and Unix domain sockets can be used: the latter allow for fast communication on a single node, whereas the former make it possible to realise a distributed computing paradigm, with clients running on different nodes or even on different HPC facilities. In order to facilitate implementation of the socket communication in client codes, a simple set of C wrappers to the standard libraries socket implementation is provided as part of the i-PI distribution, that can be used in any programming language that can be linked with C code.

As far as the communication protocol is concerned, the guiding principle has been keeping it to the lowest common denominator, and avoiding any feature that may be code-specific. Only a minimal amount of information is transferred between the client and the server; the position of the atoms and cell parameters in one direction, and the forces, virial and potential in the other.

For more details about sockets and communication, see 3.3.

1.3 Core features

i-PI includes a large number of advanced molecular dynamics features, with an obvious focus on path integral molecular dynamics, but also several methods for sampling classical trajectories.

Features in version 1.0

- molecular dynamics and PIMD in the *NVE*, *NVT* and *NPT* ensembles, with the high-frequency internal vibrations of the path propagated in the normal-mode representation [CPMM10] to allow for longer time steps;
- ring polymer contraction [MM08b][MM08a], implemented by exposing multiple socket interfaces to deal with short and long-range components of the potential energy separately; treating different components that have different computational cost and characteristic time scale separately can reduce substantially the overall effort associated with a simulation
- efficient stochastic velocity rescaling [BDP07] and path integral Langevin equation thermostats [CPMM10] to sample efficiently the canonical ensemble;
- various generalized Langevin equation (GLE) thermostats, including the optimal sampling [CBP09a][CBP10], quantum [CBP09b], and δ [CP10] thermostats; the parameters for different GLE flavors and the conditions in which they should be applied can be obtained from a separate website [Cer10];

- mixed path integral–generalized Langevin equation techniques for accelerated convergence, including both PI+GLE [CMP11] and the more recent and effective version PIGLET [CM12]; these techniques reduce the number of path integral replicas needed, while allowing for systematic convergence;
- all the standard estimators for structural properties, the quantum kinetic energy, pressure, etc.;
- more sophisticated estimators such as the scaled-coordinate heat capacity estimator [Yam05], estimators to obtain isotope fractionation free energies by re-weighting a simulation of the most abundant isotope [CM13], and a displaced-path estimator for the particle momentum distribution [LMCP10];
- the infrastructure that is needed to perform approximate quantum dynamics calculations such as ring polymer molecular dynamics (RPMD) [CM04][HMMM13] and centroid molecular dynamics (CMD) [CV93][CV94].

Further details can be found in Ref. [KRM+19].

- reweighted fourth-order path integral MD (M. Ceriotti, G.A.R. Brain) [CBRM11][JV01]; this method makes it possible to obtain fourth-order statistics by re-weighting second-order trajectories; attention should be paid to avoid statistical inefficiencies;
- finite-differences implementation of fourth-order path integrals (V. Kapil, M. Ceriotti) [KBC16]; this scheme enables explicit fourth-order path integral simulations, that converge faster than conventional Trotter methods;
- perturbed path integrals (I. Poltavsky) [PT16]; essentially, a truncated cumulant expansion of fourth-order reweighting, that often enables fast convergence avoiding statistical instability;
- open path integrals and momentum distribution estimators (V. Kapil, A. Cuzzocrea, M. Ceriotti) [KCC18]; makes it possible to compute the particle momentum distribution including quantum fluctuations of nuclei;
- quantum alchemical transformations (B. Cheng) [LAM+13][CBC16]; Monte Carlo exchanges between isotopes of different mass, useful to sample isotope propensity for different molecules or environments;
- direct isotope fractionation estimators (B. Cheng, M. Ceriotti) [CC14]; avoid thermodynamic integration to obtain isotope fractionation ratios;
- spatially localized ring polymer contraction (M. Rossi, M. Ceriotti) [LDCR18]; simple contraction scheme for weakly bound molecules, e.g. on surfaces;
- ring polymer instantons (Y. Litman, J.O. Richardson, M. Rossi); evaluation of reaction rates and tunnelling splittings for molecular rearrangements and chemical reactions;
- thermodynamic integration (M. Rossi, M. Ceriotti) [RGC16]; classical scheme to compute free energy differences;
- geometry optimizers for minimization and saddle point search (B. Helfrecht, R. Petraglia, Y. Litman, M. Rossi) [RGC16];
- harmonic vibrations through finite differences (V. Kapil, S. Bienvenue) [RGC16]; simple evaluation of the harmonic Hessian;
- multiple time stepping (V. Kapil, M. Ceriotti) [KVC16]; accelerated simulations by separating slow and fast degrees of freedom into different components of the potential energy;
- metadynamics through a PLUMED interface (G. Tribello, M. Ceriotti); simulation of rare events and free energy calculations;
- replica exchange MD (R. Petraglia, R. Meissner, M. Ceriotti) [PNicolaiW+16]; accelerated convergence of averages by performing Monte Carlo exchanges of configurations between parallel calculations
- thermostatted RPMD [RCM14], including optimized-GLE TRPMD [RKC18]; reduces well-known artifacts in the simulation of dynamical properties by path integral methods;
- dynamical corrections to Langevin trajectories (M. Rossi, V. Kapil, M. Ceriotti) [RKC18]; eliminates the artifacts introduced into dynamical properties by the presence of thermostats;

- fast forward Langevin thermostat (M. Hijazi, D. M. Wilkins, M. Ceriotti); a simple scheme to reduce the impact of strongly-damped Langevin thermostats on sampling efficiency; [HWC18]
- Langevin sampling for noisy and/or dissipative forces (J. Kessler, T. D. Kühne); suitable to stabilize and correct the artifacts that are introduced in MD trajectories by different extrapolation schemes;

1.4 Licence and credits

Most of this code is distributed under the GPL licence. For more details see www.gnu.org/licenses/gpl.html. So that they can easily be incorporated in other codes, the files in the directory “drivers” are all held under the MIT licence. For more details see <https://fedoraproject.org/wiki/Licensing:MIT>.

If you use this code in any future publications, please cite this using [CMM14] for v1 and [KRM+19] for v2.

Contributors

i-PI was originally written by M. Ceriotti and J. More at Oxford University, together with D. Manolopoulos. Several people contributed to its further development. Developers who implemented a specific feature are acknowledged above.

1.5 On-line resources

Python resources

For help with Python programming, see www.python.org. For information about the NumPy mathematical library, see www.numpy.org, and for worked examples of its capabilities see www.scipy.org/Tentative_NumPy_Tutorial. Finally, see <http://hgomersall.github.io/pyFFTW/> for documentation on the Python FFTW library that is currently implemented with i-PI.

Client code resources

Several codes provide out-of-the-box an i-PI interface, including CP2K, DFTB+, Lammmps, Quantum ESPRESSO, Siesta, FHI-aims, Yaff, deMonNano, TBE. If you are interested in interfacing your code to i-PI please get in touch, we are always glad to help!

There are several Fortran and C libraries that most client codes will probably need to run, such as FFTW, BLAS and LAPACK. These can be found at www.fftw.org, www.netlib.org/blas and www.netlib.org/lapack respectively.

These codes do not come as part of the i-PI package, and must be downloaded separately. See chapter 2.2 for more details of how to do this.

i-PI resources

For more information about i-PI and to download the source code go to <http://ipi-code.org/>.

In <http://gle4md.org/> one can also obtain colored-noise parameters to run Path Integral with Generalized Langevin Equation thermostat (PI+GLE/PIGLET) calculations.

2 Getting started

2.1 Installing i-PI

Requirements

i-PI is Python code, and as such strictly speaking does not need to be compiled and installed. The `i-pi` file in the `bin` directory of the distribution is the main (executable) script, and can be run as long as the system has installed:

- Python version 2.4 or greater (i-PI up to version v2.0 is not Python3 compatible)
- The Python numerical library NumPy

See [2.3.1](#) for more details on how to launch i-PI.

Additionally, most client codes will have their own requirements. Many of them, including the test client codes given in the “drivers” directory, will need a suitable Fortran compiler. A C compiler is required for the `sockets.c` wrapper to the `sockets` standard library. Most electronic structure codes will also need to be linked with some mathematical libraries, such as BLAS, FFTW and LAPACK. Installation instructions for these codes should be provided as part of the code distribution and on the appropriate website, as given in [1.6.2](#). Patching for use with i-PI should not introduce further dependencies.

Using the `setup.py` module

While the `i-pi` file can be used to run any i-PI simulation, it is often more convenient to install the package to the system’s Python modules path, so that it is accessible by all users and can be run without specifying the path to the Python script.

For this purpose we have included a module in the root directory of the i-PI distribution, `setup.py`, which handles creating a package with the executable and all the modules which are necessary for it to run. The first step is to build the distribution using:

```
> python setup.py build
```

Note that this requires the `distutils` package that comes with the `python-dev` package.

This creates a “build” directory containing only the files that are used to run an i-PI simulation, which can then be used to create the executable. This can be done in two ways, depending on whether or not the user has root access. If the user does have root access, then the following command will add the relevant source files to the standard Python library directory:

```
> python setup.py install
```

This will install the package in the `/usr/lib/py_vers` directory, where `py_vers` is the version of Python that is being used. This requires administrator privileges.

Otherwise, one can install i-PI in a local Python path. If such path does not exist yet, one must create directories for the package to go into, using:

Next, you must tell Python where to find this library, by appending to the Linux environment variable `PYTHONPATH`, using:

Finally, the code can be installed using:

```
> python setup.py install -prefix=
```

Either way, it will now be possible to run the code automatically, using

i-PI download

You can find information how to download i-PI from <http://ipi-code.org/download/>.

The i-PI executable, found in the `bin` folder, will run immediately, without needing to be installed or compiled, but we include a `setup.py` module in the main directory so it can be installed to the Python tree if so desired.

Installing NumPy

NumPy is the standard Python mathematics library, and is used for most of the array manipulation and linear algebra in i-PI. It should be installed alongside most standard Python environments on HPC facilities. Otherwise, it is generally relatively straightforward to install it.

In any case you must first obtain the NumPy code, which can be downloaded as a tar file from <http://www.numpy.org>. If the version of NumPy being installed is given by “`np_ver`”, this can be extracted using:

Before installing this code it first needs to be configured correctly. Note that this requires the `distutils` package that comes with the `python-dev` package. Assuming that the required software is installed, the NumPy package is built using:

```
> python setup.py build
```

The next step is to install NumPy. By default the download is to the directory `/usr/local`. If you have root access, and so can write to `/usr`, then all that needs to be done to finish the install is:

```
> python setup.py install
```

If you do not have root access, then the next step depends on which version of Python is being used. With versions 2.6 or later there is a simple command to automatically download into the directory `$HOME/local`:

```
> python setup.py install -user
```

With Python 2.4/2.5 the process is a little more involved. First you must explicitly install the package in the directory of choice, “`np_dir`” say, with the following command:

Next, you must tell Python where to find this library, by appending to the Linux environment variable `PYTHONPATH`. If you are using Python version “`py_ver`”, then the NumPy libraries will have been installed in the directory “`np_dir/lib/py_ver/site-packages`”, or a close analogue of this. In the above case the following command will allow the Python interpreter to find the NumPy libraries:

Now Python scripts can import the NumPy libraries using:

```
import numpy
```

PyFFTW

Some of the steps in the dynamics algorithm involve a change of variables from the bead coordinates to the normal modes of the ring polymers. Currently, this transformation is, at least by default, computed using a fast-Fourier transform (FFT) library within the NumPy distribution. This however is not the only distribution that could be used, and indeed faster stand-alone versions exist. The gold-standard FFT library is the FFTW library, which is a set of C libraries that have been heavily optimized for a wide range of applications. There have been a number of Python wrappers built around the FFTW library, one of which is currently interfaced with i-PI. This code can be found at <https://github.com/hgomersall/pyFFTW>, and has documentation at <http://hgomersall.github.io/pyFFTW/>.

This code has the following dependencies:

- Python version 2.7 or greater
- Numpy version 1.6 or greater
- FFTW version 3.2 or greater

This can be installed in the same way as NumPy, except using the code distribution above, or using various installation packages as per the instructions on the above documentation. Note that no other options need to be specified in the input file; i-PI will check to see if this library is available, and if it is it will be used by default. Otherwise the slower NumPy version will be used.

2.2 Installing clients

As of today, the following codes provide out-of-the-box an i-PI interface: CP2K, DFTB+, Lammmps, Quantum ESPRESSO, Siesta, FHI-aims, Yaff, deMonNano, TBE. Links to the webpages of these codes, including information on how to obtain them, can be found in <http://ipi-code.org/>.

If you are interested in interfacing your code to i-PI please get in touch, we are always glad to help. We keep some information below in case you are interested in writing a patch to a code.

Writing a patch

If you have edited a client code, and wish to make a patch available for the new version, then this can be done very simply. If your edited code is in a directory “new”, and a clean distribution is held in a directory “old”, then a patch “changes.patch” can be created using:

```
> diff -rupN old/ new/ > changes.patch
```

2.3 Running i-PI

i-PI functions based on a client-server protocol, where the evolution of the nuclear dynamics is performed by the i-PI server, whereas the energy and forces evaluation is delegated to one or more instances of an external program, that acts as a client. This design principle has several advantages, in particular the possibility of performing PIMD based on the forces produced by one’s favourite electronic structure/empirical force field code. However, it also makes running a simulation slightly more complicated, since the two components must be set up and started independently.

Running the i-PI server

i-PI simulations are run using the i-pi Python script found in the “i-pi” directory. This script takes an xml-formatted file as input, and automatically starts a simulation as specified by the data held in it. If the input file is called “input_file.xml”, then i-PI is run using:

This reads in the input data, initializes all the internally used objects, and then creates the server socket. The code will then wait until at least one client code has connected to the server before running any dynamics. Note that until this has happened the code is essentially idle, the only action that it performs is to periodically poll for incoming connections.

Running the client code

Below we give examples on how to make different clients communicate with i-PI. Most clients also include descriptions on how to do this from their own documentation.

Built-in, example client

While i-PI is designed with *ab initio* electronic structure calculations in mind, it also includes a Fortran empirical potential client code to do simple calculations and to run the examples.

The source code for this is included in the directory “drivers”, and can be compiled into an executable “i-pi-driver” using the UNIX utility make.

This code currently has four empirical potentials hardcoded into it, a Lennard-Jones potential, the Silvera-Goldman potential [SG78], a 1D harmonic oscillator potential, and the ideal gas (i.e. no potential interaction).

How the code is run is based on what command line arguments are passed to it. The command line syntax is:

```
> i-pi-driver [-u] -h hostname -p port -m [gas|lj|sg|harm] -o  
parameters [-v]
```

The flags do the following:

- u:** Optional parameter. If specified, the client will connect to a unix domain socket. If not, it will connect to an internet socket.
- h:** Is followed in the command line argument list by the hostname of the server.
- p:** Is followed in the command line argument list by the port number of the server.
- m:** Is followed in the command line argument list by a string specifying the type of potential to be used. “gas” gives no potential, “lj” gives a Lennard-Jones potential, “sg” gives a Silvera-Goldman potential and “harm” gives a 1D harmonic oscillator potential. Other options should be clear from their description.
- o:** Is followed in the command line argument list by a string of comma separated values needed to initialize the potential parameters. “gas” requires no parameters, “harm” requires a spring constant, “sg” requires a cut-off radius and “lj” requires the length and energy scales and a cut-off radius to be specified. All of these must be given in atomic units.
- v:** Optional parameter. If given, the client will print out more information each time step.

This code should be fairly simple to extend to other pair-wise interaction potentials, and examples of its use can be seen in the “examples” directory, as explained in 2.4.

CP2K

To use CP2K as the client code using an internet domain socket on the host address “host_address” and on the port number “port” the following lines must be added to its input file:

If instead a unix domain socket is required then the following modification is necessary:

The rest of the input file should be the same as for a standard CP2K calculation, as explained at www.cp2k.org/.

Quantum-Espresso

To use Quantum-Espresso as the client code using an internet domain socket on the host address “host_address” and on the port number “port” the following lines must be added to its input file:

If instead a unix domain socket is required then the following modification is necessary:

The rest of the input file should be the same as for a standard Quantum Espresso calculation, as explained at www.quantum-espresso.org/.

LAMMPS

To use LAMMPS as the client code using an internet domain socket on the host address “host_address” and on the port number “port” the following lines must be added to its input file:

If instead a unix domain socket is required then the following modification is necessary:

The rest of the input file should be the same as for a standard LAMMPS calculation, as explained at <http://lammps.sandia.gov/index.html>. Note that LAMMPS must be compiled with the `yes-user-misc` option to communicate with i-PI. More information from https://lammps.sandia.gov/doc/fix_ipi.html.

FHI-aims

To use FHI-aims as the client code using an internet domain socket on the host address “host_address” and on the port number “port” the following lines must be added to its `control.in` file:

If instead a unix domain socket is required then the following modification is necessary:

One can also communicate different electronic-structure quantities to i-PI through the `extra` string from FHI-aims. In this case the following lines can be added to the `control.in` file:

where option can be, e.g., `dipole`, `hirshfeld`, `workfunction`, `friction`.

Running on a HPC system

Running i-PI on a high-performance computing (HPC) system can be a bit more challenging than running it locally using UNIX-domain sockets or using the `localhost` network interface. The main problem is related to the fact that different HPC systems adopt a variety of solutions to have the different nodes communicate with each other and with the login nodes, and to queue and manage computational jobs.

Figure 2.1 represents schematically three different approaches to run i-PI on a HPC system:

1. running both i-PI and multiple instances of the client as a single job on the HPC system. The job submission script must launch i-PI first, as a serial background job, then wait a few seconds for it to load and create a socket. Then, one should launch with `mpirun` or any system-specific mechanism one or more independent instances of the client code. Note that not all queuing systems allow launching several `mpirun` instances from a single job.




Fig. 4: Different approaches to run i-PI and a number of instances of the forces code on a HPC system: a) running i-PI and the clients in a single job; b) running i-PI and the clients on the same system, but using different jobs, or running i-PI interactively on the login node; c) running i-PI on a local workstation, communicating with the clients (that can run on one or multiple HPC systems) over the internet.

2. running i-PI and the clients on the HPC system, but in separate jobs. Since i-PI consumes very little resources, one should ideally launch it interactively on a login node

or alternative on a queue with a very long wall-clock time. Then, multiple instances of the client can be run as independent jobs: as they start, they will connect to the server which will take care of adding them dynamically to the list of active clients, dispatching force calculations to them, and removing them from the list when their wall-clock time expires. This is perhaps the model that applies more easily to different HPC systems; however it requires having permission to run on the head node, or having access to a long wall-clock time queue that ensures that i-PI is always active.

3. running i-PI on a simple workstation, and performing communication over the internet with the clients that run on one or more HPC systems. This model exploits in full the distributed-computing model that underlies the philosophy of i-PI and is very robust – as the server can be always on, and the output of the simulation is generated locally. However, this is also the most complicated to set up, as the local workstation must accept incoming connections from the internet – which is not always possible when behind a firewall – and the compute nodes of the HPC centre must have an outgoing connection to the internet, which often requires ssh tunnelling through a login node (see section 3.3 for more details).

2.4 Testing the install

test the installation with ‘nose’

There are several test cases included, that can be run automatically with ‘i-pi-tests’ from the root directory.

```
> i-pi-tests
```


test cases and examples

The *examples/* folder contain a multitude of examples for i-PI, covering most of the existing functionalities, and including also simple tests that can be run with different client codes.

All the input files are contained in the directory “examples”, which is subdivided into subfolder that cover different classes of simulations, and/or different client codes.

3 User guide

3.1 Input files

Input file format and structure

In order to give the clearest layout, xml formatting was chosen as the basis for the main input file. An xml file consists of a set of hierarchically nested tags. There are three parts to an xml tag. Each tag is identified by a tag name, which specifies the class or variable that is being initialized. Between the opening and closing tags there may be some data, which may or may not contain other tags. This is used to specify the contents of a class object, or the value of a variable. Finally tags can have attributes, which are used for metadata, i.e. data used to specify how the tag should be interpreted. As an example, a ‘mode’ attribute can be used to select between different thermostating algorithms, specifying how the options of the thermostat class should be interpreted.

A xml tag has the following syntax:

The syntax for the different types of tag data is given below:

Data type	Syntax
Boolean	<tag>True</tag> or <tag>False</tag>
Float	<tag>11.111</tag> or <tag>1.111e+1</tag>
Integer	<tag>12345</tag>
String	<tag>string_data</tag>
Tuple	<tag> (int1, int2, ...) </tag>
Array	<tag> [entry1, entry2, ...] </tag>
Dictionary	<tag>{name1: data1, name2: data2, ... } </tag>

Note that arrays are always given as one-dimensional lists. In cases where a multi-dimensional array must be entered, one can use the ‘shape’ attribute, that determines how the list will be reshaped into a multi-dimensional array. For example, the bead positions are represented in the code as an array of shape (number of beads, 3*number of atoms). If we take a system with 20 atoms and 8 beads, then this can be represented in the xml input file as:

```
<beads nbeads='8' natoms='20'> <q shape='(8,60)'+[ q11x, q11y, q11z,  
q12x, q12y, ... ]</q> ... </beads>
```

If ‘shape’ is not specified, a 1D array will be assumed.

The code uses the hierarchical nature of the xml format to help read the data; if a particular object is held within a parent object in the code, then the tag for that object will be within the appropriate parent tags. This is used to make the structure of the simulation clear.

For example, the system that is being studied is partly defined by the thermodynamic ensemble that should be sampled, which in turn may be partly defined by the pressure, and so on. To make this dependence clear in the code the global simulation object which holds all the data contains an ensemble object, which contains a pressure variable.

Therefore the input file is specified by having a tag, containing an tag, which itself contains a “pressure” tag, which will contain a float value corresponding to the external pressure. In this manner, the class structure can be constructed iteratively.

For example, suppose we want to generate a *NPT* ensemble at an external pressure of 10^{-7} atomic pressure units. This would be specified by the following input file:

```
<system> <ensemble> <pressure> 1e-7 </pressure> ... </ensemble> ...  
</system>
```

To help detect any user error the recognized tag names, data types and acceptable options are all specified in the code in a specialized input class for each class of object. A full list of all the available tags and a brief description of their function is given in chapter 4.

Overriding default units

Many of the input parameters, such as the pressure in the above example, can be specified in more than one unit. Indeed, often the atomic unit is inconvenient to use, and we would prefer something else. Let us take the above example, but instead take an external pressure of 3 MPa. Instead of converting this to the atomic unit of pressure, it is possible to use pascals directly using:

```
<system> <ensemble> <pressure units='pascal'> 3e6 </pressure> ...  
</ensemble> ... </system>
```

The code can also understand S.I. prefixes, so this can be simplified further using:

```
<system> <ensemble> <pressure units='megapascal'> 3 </pressure> ...  
</ensemble> ... </system>
```

A full list of which units are defined for which dimensions can be found in the `units.py` module.

Initialization section

The input file can contain a tag, which contains a number of fields that determine the starting values of the various quantities that define the state of the simulation – atomic positions, cell parameters, velocities, These fields (, , , ,) specify how the values should be obtained: either from a manually-input list or from an external file.

Configuration files

Instead of initializing the atom positions manually, the starting configuration can be specified through a separate data file. The name of the configuration file is specified within one of the possible fields of an tag. The file format is specified with the “mode” attribute. The currently accepted file formats are:

- `pdb`
- `xyz`
- `chk`

the last of which will be described in the next section.

Depending on the field name, the values read from the external file will be used to initialize one component of the simulation or another (e.g. the positions or the velocities). The tag can be used as a shortcut to initialize the atom positions, labels, masses and possibly the cell parameters at the same time. For instance,

```
<initialize nbeads="8"> <file mode="pdb"> init.pdb </file>
</initialize>
```

is equivalent to

```
<initialize nbeads="8"> <positions mode="pdb"> init.pdb </positions>
<labels mode="pdb"> init.pdb </labels> <masses mode="pdb"> init.pdb
</masses> <cell mode="pdb"> init.pdb </cell> </initialize>
```

In practice, the using the tag will only read the information that can be inferred from the given file type, so for an ‘xyz’ file, the cell parameters will not be initialized.

Initialization from checkpoint files

i-PI gives the option to output the entire state of the simulation at a particular timestep as an xml input file, called a checkpoint file (see 3.2.3 for details). As well as being a valid input for i-PI, a checkpoint can also be used inside an tag to specify the configuration of the system, discarding other parameters of the simulation such as the current time step or the chosen ensemble. Input from a checkpoint is selected by using “chk” as the value of the “mode” attribute. As for the configuration file, a checkpoint file can be used to initialize either one or many variables depending on which tag name is used.

3.2 Output files

i-PI uses a very flexible mechanism to specify how and how often atomic configurations and physical properties should be output. Within the tag of the xml input file the user can specify multiple tags, each one of which will correspond to a particular output file. Each file is managed separately by the code, so what is output to a particular file and how often can be adjusted for different files independently.

For example, some of the possible output properties require more than one force evaluation per time step to calculate, and so can considerably increase the computational cost of a simulation unless they are computed once every several time steps. On the other hand, for properties such as the conserved energy quantity it is easy, and often useful, to output them every time step as they are simple to compute and do not take long to output to file.

There are three types of output file that can be specified; property files for system level properties, trajectory files for atom/bead level properties, and checkpoint files which save the state of the system and so can be used to restart the simulation from a particular point. For a brief overview of the format of each of these types of files, and some of their more common uses, see 5.1. To give a more in depth explanation of each of these files, they will now be considered in turn.

Properties

This is the output file for all the system and simulation level properties, such as the total energy and the time elapsed. It is designed to track a small number of important properties throughout a simulation run, and as such has been formatted to be used as input for plotting programs such as gnuplot.

The file starts with a header, which describes the properties being written in the different columns and their output units. This is followed by the actual data. Each line corresponds to one instant of the simulation. The file is fixed formatted, with two blank characters at the start of each row, then the data in the same order as the header row. By default, each column is 16 characters wide and every float is written in exponential format with 8 digits after the decimal point.

For example, if we had asked for the current time step, the total simulation time in picoseconds, and the potential energy in electronvolt, then the properties output file would look something like:

```
# column 1 -> step : The current simulation time step. # column 2 ->
timepicosecond : The elapsed simulation time. # column 3 ->
potentialelectronvolt : The physical system potential energy.
0.00000000e+00 0.00000000e+00 -1.32860475e+04 1.00000000e+00
1.00000000e-03 -1.32865789e+04 ...
```

The properties that are output are determined by the tag in the xml input file. The format of this tag is:

```
<properties stride= filename= flush= shape=> [ prop1nameunits(arg1;
... ), prop2name...(.), ... ] </properties>
```

e.g.

The attributes have the following meanings:

stride

The number of steps between each output to file

filename

The name of the output file

flush

The number of output lines between buffer flushes

shape

The number of properties in the list.

The tag data is an array of strings, each of which contains three different parts:

- The property name, which describes which type of property is to be output. This is a mandatory part of the string.
- The units that the property will be output in. These are specified between curly brackets. If this is not specified, then the property will be output in atomic units. Note that some properties can only be output in atomic units.
- The arguments to be passed to the function. These are specified between standard brackets, with each argument separated by a semi-colon. These may or may not be mandatory depending on the property, as some arguments have well defined default values. The arguments can be specified by either of two different syntaxes, (name1=arg1; ...) or (arg1; ...).

The first syntax uses keyword arguments. The above example would set the variable with the name “name1” the value “arg1”. The second syntax uses positional arguments. This syntax relies on the arguments being specified in the correct order, as defined in the relevant function in the `properties.py` module, since the user has not specified which variable to assign the value to.

The two syntaxes may be mixed, but positional arguments must be specified first otherwise undefined behaviour will result. If no arguments are specified, then the defaults as defined in the `properties.py` module will be used.

List of available properties

atom_f: The force (x,y,z) acting on a particle given its index. Takes arguments index and bead (both zero based). If bead is not specified, refers to the centroid.

dimension: force; size: 3;

atom_f_path: The forces acting on all the beads of a particle given its index. Takes arguments index and bead (both zero based). If bead is not specified, refers to the centroid.

dimension: length;

atom_p: The momentum (x,y,z) of a particle given its index. Takes arguments index and bead (both zero based). If bead is not specified, refers to the centroid.

dimension: momentum; size: 3;

atom_v: The velocity (x,y,z) of a particle given its index. Takes arguments index and bead (both zero based). If bead is not specified, refers to the centroid.

dimension: velocity; size: 3;

atom_x: The position (x,y,z) of a particle given its index. Takes arguments index and bead (both zero based). If bead is not specified, refers to the centroid.

dimension: length; size: 3;

atom_x_path: The positions of all the beads of a particle given its index. Takes an argument index (zero based).

dimension: length;

bead_potentials: The physical system potential energy of each bead.

dimension: energy; size: nbeads;

bweights_component: The weight associated one part of the hamiltonian. Takes one mandatory argument index (zero-based) that indicates for which component of the hamiltonian the weight must be returned.

cell_abcABC: The lengths of the cell vectors and the angles between them in degrees as a list of the form [a, b, c, A, B, C], where A is the angle between the sides of length b and c in degrees, and B and C are defined similarly. Since the output mixes different units, a, b and c can only be output in bohr.

size: 6;

cell_h: The simulation cell as a matrix. Returns the 6 non-zero components in the form [xx, yy, zz, xy, xz, yz].

dimension: length; size: 6;

chin_weight: The 3 numbers output are 1) the logarithm of the weighting factor $-\beta_P \Delta H$, 2) the square of the logarithm, and 3) the weighting factor

size: 3;

conserved: The value of the conserved energy quantity per bead.

dimension: energy;

density: The mass density of the physical system.

dimension: density;

displacedpath: This is the estimator for the end-to-end distribution, that can be used to calculate the particle momentum distribution as described in L. Lin, J. A. Morrone, R. Car and M. Parrinello, 105, 110602 (2010), Phys. Rev. Lett. Takes arguments 'ux', 'uy' and 'uz', which are the components of the path opening vector. Also takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the end-to-end distribution estimator for. If not specified, all atoms are used. Note that one atom is computed at a time, and that each path opening operation costs as much as a PIMD step. Returns the average over the selected atoms of the estimator of $\exp(-U(u))$ for each frame.

ensemble_bias: The bias applied to the current ensemble

dimension: energy;

ensemble_lp: The log of the ensemble probability

ensemble_pressure: The target pressure for the current ensemble

dimension: pressure;

ensemble_temperature: The target temperature for the current ensemble

dimension: temperature;

forcemod: The modulus of the force. With the optional argument 'bead' will print the force associated with the specified bead.

dimension: force;

hweights_component: The weight associated one part of the hamiltonian. Takes one mandatory argument index (zero-based) that indicates for which component of the hamiltonian the weight must be returned.

isotope_scsep: Returns the (many) terms needed to compute the scaled-coordinates free energy perturbation scaled mass KE estimator (M. Ceriotti, T. Markland, J. Chem. Phys. 138, 014112 (2013)). Takes two arguments, 'alpha' and 'atom', which give the scaled mass parameter and the atom of interest respectively, and default to '1.0' and '.'. The 'atom' argument can either be the label of a particular kind of atom, or an index (zero based) of a specific atom. This property computes, for each atom in the selection, an estimator for the kinetic energy it would have had if it had the mass scaled by alpha. The 7 numbers output are the average over the selected atoms of the log of the weights $\langle h \rangle$, the average of the squares $\langle h^2 \rangle$, the average of the un-weighted scaled-coordinates kinetic energies $\langle T_{CV} \rangle$ and of the squares $\langle T_{CV}^2 \rangle$, the log sum of the weights $LW = \ln(\sum(e^{-(h)}))$, the sum of the re-weighted kinetic energies, stored as a log modulus and sign, $LTW = \ln(\text{abs}(\sum(T_{CV} e^{-(h)})))$ $STW = \text{sign}(\sum(T_{CV} e^{-(h)}))$. In practice, the best estimate of the estimator can be computed as $[\sum_i \exp(LTW_i) * STW_i] / [\sum_i \exp(LW_i)]$. The other terms can be used to compute diagnostics for the statistical accuracy of the re-weighting process. Note that evaluating this estimator costs as much as a PIMD step for each atom in the list. The elements that are output have different units, so the output can be only in atomic units.

size: 7;

isotope_tdfep: Returns the (many) terms needed to compute the thermodynamic free energy perturbation scaled mass KE estimator (M. Ceriotti, T. Markland, J. Chem. Phys. 138, 014112 (2013)). Takes two arguments, 'alpha' and 'atom', which give the scaled mass parameter and the atom of interest respectively, and default to '1.0' and '.'. The 'atom' argument can either be the label of a particular kind of atom, or an index (zero based) of a specific atom. This property computes, for each atom in the selection, an estimator for the kinetic energy it would have had if it had the mass scaled by alpha. The 7 numbers output are the average over the selected atoms of the log of the weights $\langle h \rangle$, the average of the squares $\langle h^2 \rangle$, the average of the un-weighted scaled-coordinates kinetic energies $\langle T_{CV} \rangle$ and of the squares $\langle T_{CV}^2 \rangle$, the log sum of the weights $LW = \ln(\sum(e^{-(h)}))$, the sum of the re-weighted kinetic energies, stored as a log modulus and sign, $LTW = \ln(\text{abs}(\sum(T_{CV} e^{-(h)})))$ $STW = \text{sign}(\sum(T_{CV} e^{-(h)}))$. In practice, the best estimate of the estimator can be computed as $[\sum_i \exp(LTW_i) * STW_i] / [\sum_i \exp(LW_i)]$. The other terms can be used to compute diagnostics for the statistical accuracy of the re-weighting process. Evaluating this estimator is inexpensive, but typically the statistical accuracy is worse than with the scaled coordinates estimator. The elements that are output have different units, so the output can be only in atomic units.

size: 7;

isotope_zetasc: Returns the (many) terms needed to directly compute the relative probability of isotope substitution in two different systems/phases. Takes four arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. The 3 numbers output are 1) the average over the excess potential energy for scaled coordinates $\langle sc \rangle$, 2) the average of the squares of the excess potential energy $\langle sc^2 \rangle$, and 3) the average of the exponential of excess potential energy $\langle \exp(-\beta * sc) \rangle$

size: 3;

isotope_zetasc_4th: Returns the (many) terms needed to compute the scaled-coordinates fourth-order direct estimator. Takes two arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. The 5 numbers output are 1) the average over the excess potential energy for an isotope atom substitution $\langle sc \rangle$, 2) the average of the squares of the excess potential energy $\langle sc^2 \rangle$, and 3) the average of the exponential of excess potential energy $\langle \exp(-\beta * sc) \rangle$, and 4-5) Suzuki-Chin and Takahashi-Imada 4th-order reweighing term

size: 5;

isotope_zetatd: Returns the (many) terms needed to directly compute the relative probability of isotope substitution in two different systems/phases. Takes two arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. The 3 numbers output are 1) the average over the excess spring energy for an isotope atom substitution $\langle \text{spr} \rangle$, 2) the average of the squares of the excess spring energy $\langle \text{spr}^2 \rangle$, and 3) the average of the exponential of excess spring energy $\langle \exp(-\beta \text{spr}) \rangle$

size: 3;

isotope_zetatd_4th: Returns the (many) terms needed to compute the thermodynamic fourth-order direct estimator. Takes two arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. The 5 numbers output are 1) the average over the excess spring energy for an isotope atom substitution $\langle \text{spr} \rangle$, 2) the average of the squares of the excess spring energy $\langle \text{spr}^2 \rangle$, and 3) the average of the exponential of excess spring energy $\langle \exp(-\beta \text{spr}) \rangle$, and 4-5) Suzuki-Chin and Takahashi-Imada 4th-order reweighing term

size: 5;

kcv_scaledcoords: Returns the estimators that are required to evaluate the scaled-coordinates estimators for total energy and heat capacity, as described in T. M. Yamamoto, J. Chem. Phys., 104101, 123 (2005). Returns eps_v and eps_v' , as defined in that paper. As the two estimators have a different dimensions, this can only be output in atomic units. Takes one argument, 'fd_delta', which gives the value of the finite difference parameter used - which defaults to -0.0001. If the value of 'fd_delta' is negative, then its magnitude will be reduced automatically by the code if the finite difference error becomes too large.

size: 2;

kinetic_cv: The centroid-virial quantum kinetic energy of the physical system. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the kinetic energy of. If not specified, all atoms are used.

dimension: energy;

kinetic_ij: The centroid-virial off-diagonal quantum kinetic energy tensor of the physical system. This computes the cross terms between atoms i and atom j, whose average is $\langle p_i p_j / (2 \sqrt{m_i m_j}) \rangle$. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz]. Takes arguments 'i' and 'j', which give the indices of the two desired atoms.

dimension: energy; size: 6;

kinetic_md: The kinetic energy of the (extended) classical system. Takes optional arguments 'atom', 'bead' or 'nm'. 'atom' can be either an atom label or an index (zero-based) to specify which species or individual atom to output the kinetic energy of. If not specified, all atoms are used and averaged. 'bead' or 'nm' specify whether the kinetic energy should be computed for a single bead or normal mode. If not specified, all atoms/beads/nm are used.

dimension: energy;

kinetic_opsc: The centroid-virial quantum kinetic energy of the physical system. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the kinetic energy of. If not specified, all atoms are used.

dimension: energy;

kinetic_prsc: The Suzuki-Chin primitive estimator of the quantum kinetic energy of the physical system

dimension: energy;

kinetic_td: The primitive quantum kinetic energy of the physical system. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the kinetic energy of. If not specified, all atoms are used.

dimension: energy;

kinetic_tdsc: The Suzuki-Chin centroid-virial thermodynamic estimator of the quantum kinetic energy of the physical system. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the kinetic energy of. If not specified, all atoms are used.

dimension: energy;

kinetic_tens: The centroid-virial quantum kinetic energy tensor of the physical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz]. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the kinetic tensor components of. If not specified, all atoms are used.

dimension: energy; size: 6;

kstress_cv: The quantum estimator for the kinetic stress tensor of the physical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

kstress_md: The kinetic stress tensor of the (extended) classical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

kstress_tdsc: The Suzuki-Chin thermodynamic estimator for pressure of the physical system.

dimension: pressure;

pot_component: The contribution to the system potential from one of the force components. Takes one mandatory argument index (zero-based) that indicates which component of the potential must be returned. The optional argument 'bead' will print the potential associated with the specified bead. If the potential is weighed, the weight will be applied.

dimension: energy;

pot_component_raw: The contribution to the system potential from one of the force components. Takes one mandatory argument index (zero-based) that indicates which component of the potential must be returned. The optional argument 'bead' will print the potential associated with the specified bead. Potential weights will not be applied.

dimension: energy;

potential: The physical system potential energy. With the optional argument 'bead' will print the potential associated with the specified bead.

dimension: energy;

potential_opsc: The Suzuki-Chin operator estimator for the potential energy of the physical system.

dimension: energy;

potential_tdsc: The Suzuki-chin thermodynamic estimator for the potential energy of the physical system.

dimension: energy;

pressure_cv: The quantum estimator for pressure of the physical system.

dimension: pressure;

pressure_md: The pressure of the (extended) classical system.

dimension: pressure;

pressure_tdsc: The Suzuki-Chin thermodynamic estimator for pressure of the physical system.

dimension: pressure;

r_gyration: The average radius of gyration of the selected ring polymers. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the radius of gyration of. If not specified, all atoms are used and averaged.

dimension: length;

sc_scaledcoords: Returns the estimators that are required to evaluate the scaled-coordinates estimators for total energy and heat capacity, as described in T. M. Yamamoto, J. Chem. Phys., 104101, 123 (2005). Returns `eps_v` and `eps_v'`, as defined in that paper. As the two estimators have a different dimensions, this can only be output in atomic units. Takes one argument, 'fd_delta', which gives the value of the finite difference parameter used - which defaults to -0.0001. If the value of 'fd_delta' is negative, then its magnitude will be reduced automatically by the code if the finite difference error becomes too large.

size: 2;

scaledcoords: Returns the estimators that are required to evaluate the scaled-coordinates estimators for total energy and heat capacity, as described in T. M. Yamamoto, J. Chem. Phys., 104101, 123 (2005). Returns `eps_v` and `eps_v'`, as defined in that paper. As the two estimators have a different dimensions, this can only be output in atomic units. Takes one argument, 'fd_delta', which gives the value of the finite difference parameter used - which defaults to -0.0001. If the value of 'fd_delta' is negative, then its magnitude will be reduced automatically by the code if the finite difference error becomes too large.

size: 2;

spring: The total spring potential energy between the beads of all the ring polymers in the system.

dimension: energy;

step: The current simulation time step.

dimension: number;

stress_cv: The total quantum estimator for the stress tensor of the physical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

stress_md: The total stress tensor of the (extended) classical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

temperature: The current temperature, as obtained from the MD kinetic energy of the (extended) ring polymer. Takes optional arguments 'atom', 'bead' or 'nm'. 'atom' can be either an atom label or an index (zero-based) to specify which species or individual atom to output the temperature of. If not specified, all atoms are used and averaged. 'bead' or 'nm' specify whether the temperature should be computed for a single bead or normal mode.

dimension: temperature;

ti_pot: The correction potential in Takahashi-Imada 4th-order PI expansion. Takes an argument 'atom', which can be either an atom label or index (zero based) to specify which species to find the correction term for. If not specified, all atoms are used.

dimension: energy; size: 1;

ti_weight: The 3 numbers output are 1) the logarithm of the weighting factor $-\beta_P \Delta H$, 2) the square of the logarithm, and 3) the weighting factor

size: 3;

time: The elapsed simulation time.

dimension: time;

vcom: The center of mass velocity (x,y,z) of the system or of a species. Takes arguments label (default to all species) and bead (zero based). If bead is not specified, refers to the centroid.

dimension: velocity; size: 3;

vir_tdsc: The Suzuki-Chin thermodynamic estimator for pressure of the physical system.

dimension: pressure;

virial_cv: The quantum estimator for the virial stress tensor of the physical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

virial_fq: Returns the scalar product of force and positions. Useful to compensate for the harmonic component of a potential. Gets one argument ‘ref’ that should be a filename for a reference configuration, in the style of the FFDebye geometry input, and one that contains the input units.

dimension: energy; size: 1;

virial_md: The virial tensor of the (extended) classical system. Returns the 6 independent components in the form [xx, yy, zz, xy, xz, yz].

dimension: pressure; size: 6;

volume: The volume of the cell box.

dimension: volume;

Trajectory files

These are the output files for atomic or bead level properties, such as the bead positions. In contrast to properties files, they output data for all atomic degrees of freedom, in a format that can be read by visualization packages such as VMD.

Multiple trajectory files can be specified, each described by a separate tag within the section of the input file. The allowable file formats for the trajectory output files are the same as for the configuration input files, given in [3.1.2.1](#).

These tags have the format:

This is very similar to the tag, except that it has the additional tags “format” and “cell_units”, and only one *traj_name* quantity can be specified per file. ‘format’ specifies the format of the output file, and ‘cell_units’ specifies the units in which the cell dimensions are output. Depending on the quantity being output, the trajectory may consist of just one file per time step (e.g. the position of the centroid) or of several files, one for each bead, whose name will be automatically determined by appending the bead index to the specified “filename” attribute (e.g. the beads position). In the latter case it is also possible to output the quantity computed for a single bead by specifying its (zero-based) index in the “bead” attribute.

List of available trajectory files

extras: The additional data returned by the client code. If the attribute “extra_type” is specified, and if the data is JSON formatted, it prints only the specified field. Otherwise (or if extra_type=“raw”) the full string is printed verbatim. Will print out one file per bead, unless the bead attribute is set by the user.

extras_component: The additional data returned by the client code, printed verbatim or expanded as a dictionary. See “extras”. Fetches the extras from a specific force component, indicated in parentheses [extras_component(idx)].

f_centroid: The force acting on the centroid.

dimension: force;

forces: The force trajectories. Will print out one file per bead, unless the bead attribute is set by the user.

dimension: force;

forces_sc: The Suzuki-Chin component of force trajectories. Will print out one file per bead, unless the bead attribute is set by the user.

dimension: force;

isotope_zetasc: Scaled-coordinates isotope fractionation direct estimator in the form of ratios of partition functions. Takes two arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. All the atoms but the selected ones will have zero output

isotope_zetatd: Thermodynamic isotope fractionation direct estimator in the form of ratios of partition functions. Takes two arguments, 'alpha', which gives the scaled mass parameter and default to '1.0', and 'atom', which is the label or index of a type of atoms. All the atoms but the selected ones will have zero output

kinetic_cv: The centroid virial quantum kinetic energy estimator for each atom, resolved into Cartesian components [xx, yy, zz]

dimension: energy;

kinetic_od: The off diagonal elements of the centroid virial quantum kinetic energy tensor [xy, xz, yz]

dimension: energy;

momenta: The momentum trajectories. Will print out one file per bead, unless the bead attribute is set by the user.

dimension: momentum;

p_centroid: The centroid momentum.

dimension: momentum;

positions: The atomic coordinate trajectories. Will print out one file per bead, unless the bead attribute is set by the user.

dimension: length;

r_gyration: The radius of gyration of the ring polymer, for each atom and resolved into Cartesian components [xx, yy, zz]

dimension: length;

v_centroid: The centroid velocity.

dimension: velocity;

v_centroid_even: The suzuki-chin centroid velocity.

dimension: velocity;

v_centroid_odd: The suzuki-chin centroid velocity.

dimension: velocity;

velocities: The velocity trajectories. Will print out one file per bead, unless the bead attribute is set by the user.

dimension: velocity;

x_centroid: The centroid coordinates.

dimension: length;

x_centroid_even: The suzuki-chin centroid coordinates.

dimension: length;

x_centroid_odd: The suzuki-chin centroid coordinates.

dimension: length;

Checkpoint files

As well as the above output files, the state of the system at a particular time step can also be saved to file. These checkpoint files can later be used as input files, with all the information required to restore the state of the system to the point at which the file was created.

This is specified by the tag which has the syntax:

```
<checkpoint stride= filename= overwrite=> step </checkpoint>
```

Again, this is similar to the `and` tags, but instead of having a value which specifies what to output, the value simply gives a number to identify the current checkpoint file. There is also one additional attribute, “`overwrite`”, which specifies whether each new checkpoint file overwrites the old one, or whether all checkpoint files are kept. If they are kept, they will be written not to the file “`filename`”, but instead an index based on the value of “`step`” will be appended to it to distinguish between different files.

If the ‘`step`’ parameter is not specified, the following syntax can also be used:

```
<checkpoint stride= filename= overwrite=/>
```

Soft exit and RESTART

As well as outputting checkpoint files during a simulation run, i-PI also creates a checkpoint automatically at the end of the simulation, with file name “`RESTART`”. In the same way as the checkpoint files discussed above, it contains the full state of the simulation. It can be used to seamlessly restart the simulation if the user decides that a longer run is needed to gather sufficient statistics, or if i-PI is terminated before the desired number of steps have been completed.

i-PI will try to generate a `RESTART` file when it terminates, either because *total_time* has elapsed, or because it received a (soft) kill signal by the operating system. A soft exit can also be forced by creating an empty file named “`EXIT`” in the directory in which i-PI is running.

An important point to note is that since each time step is split into several parts, it is only at the end of each step that all the variables are consistent with each other in such a way that the simulation can be restarted from them without changing the dynamics. Thus if a soft exit call is made during a step, then the restart file that is created must correspond to the state of the system *before* that step began. To this end, the state of the system is saved at the start of every step.

In order to restart i-PI from a file named “`RESTART`” one simply has to run

```
> python i-pi RESTART
```

3.3 Distributed execution

Communication protocol

i-PI is based on a clear-cut separation between the evolution of the nuclear coordinates and the evaluation of energy and forces, which is delegated to an external program. The two parts are kept as independent as possible, to minimize the client-side implementation burden, and to make sure that the server will be compatible with any empirical or *ab initio* code that can compute inter-atomic forces for a given configuration.

Once a communication channel has been established between the client and the server (see 3.3.3), the two parties exchange minimal information: i-PI sends the atomic positions and the cell parameters to the client, which computes energy, forces and virial and returns them to the server.

The exchange of information is regulated by a simple communication protocol. The server polls the status of the client, and when the client signals that it is ready to compute forces i-PI sends the atomic positions to it. When the client

responds to the status query by signalling that the force evaluation is finished, i-PI will prepare to receive the results of the calculation. If at any stage the client does not respond to a query, the server will wait and try again until a prescribed timeout period has elapsed, then consider the client to be stuck, disconnect from it and reassign its force evaluation task to another active instance. The server assumes that 4-byte integers, 8-byte floats and 1-byte characters are used. The typical communication flow is as follows:

1. a header string “**STATUS**” is sent by the server to the client that has connected to it;
2. a header string is then returned, giving the status of the client code. Recognized messages are:

“NEEDINIT”:

if the client code needs any initialising data, it can be sent here. The server code will then send a header string “INIT”, followed by an integer corresponding to the bead index, another integer giving the number of bits in the initialization string, and finally the initialization string itself.

“READY”:

sent if the client code is ready to calculate the forces. The server socket will then send a string “POSDATA”, then nine floats for the cell vector matrix, then another nine floats for the inverse matrix. The server socket will then send one integer giving the number of atoms, then the position data as 3 floats for each atom giving the 3 cartesian components of its position.

“HAVEDATA”:

is sent if the client has finished computing the potential and forces. The server socket then sends a string “GETFORCE”, and the client socket returns “FORCEREADY”. The potential is then returned as a float, the number of atoms as an integer, then the force data as 3 floats per atom in the same way as the positions, and the virial as 9 floats in the same way as the cell vector matrix. Finally, the client may return an arbitrary string containing additional data that have been obtained by the electronic structure calculation (atomic charges, dipole moment, ...). The client first returns an integer specifying the number of characters, and then the string, which will be output verbatim if this “extra” information is requested in the output section (see 3.2.2). The string can be formatted in the JSON format, in which case i-PI can extract and process individual fields, that can be printed separately to different files.

3. The server socket waits until the force data for each replica of the system has been calculated and returned, then the MD can be propagated for one more time step, and new force requests will be dispatched.

Parallelization

As mentioned before, one of the primary advantages of using this type of data transfer is that it allows multiple clients to connect to an i-PI server, so that different replicas of the system can be assigned to different client codes and their forces computed in parallel. In the case of *ab initio* force evaluation, this is a trivial level of parallelism, since the cost of the force calculation is overwhelming relative to the overhead involved in exchanging coordinates and forces. Note that even if the parallelization over the replicas is trivial, often one does not obtain perfect scaling, due to the fact that some of the atomic configurations might require more steps to reach self-consistency, and the wall-clock time per step is determined by the slowest replica.

i-PI maintains a list of active clients, and distributes the forces evaluations among those available. This means that, if desired, one can run an n -bead calculation using only $m < n$ clients, as the server takes care of sending multiple replicas to each client per MD step. To avoid having clients idling for a substantial amount of time, m should be a divisor of n . The main advantage of this approach, compared to one that rigidly assigns one instance of the client to each bead, is that if each client is run as an independent job in a queue (see 2.3.3), i-PI can start performing PIMD as soon as a single job has started, and can carry on advancing the simulation even if one of the clients becomes unresponsive.

Especially for *ab initio* calculations, there is an advantage in running with $m = n$. i-PI will always try to send the coordinates for one path integral replica to the client that computed it at the previous step: this reduces the change in the particle positions between force evaluations, so that the charge density/wavefunction from the previous step is a better starting guess and self-consistency can be achieved faster. Also, receiving coordinates that represent a continuous trajectory makes it possible to use extrapolation strategies that might be available in the client code.

Obviously, most electronic-structure client codes provide a further level of parallelisation, based on OpenMP and/or MPI. This is fully compatible with i-PI, as it does not matter how the client does the calculation since only the forces, potential and virial are sent to the server, and the communication is typically performed by the master process of the client.

Sockets

The communication between the i-PI server and the client code that evaluates forces is implemented through sockets. A socket is a data transfer device that is designed for internet communication, so it supports both multiple client connections to the same server and two-way communication. This makes sockets ideal for use in i-PI, where each calculation may require multiple instances of the client code. A socket interface can actually function in two different modes.

UNIX-domain sockets are a mechanism for local, inter-process communication. They are fast, and best suited when one wants to run i-PI with empirical potentials, and the latency of the communication with the client becomes a significant overhead for the calculation. UNIX-domain sockets create a special file in the local file system, that serves as a rendezvous point between server and clients, and are uniquely identified by the name of the file itself, that can be specified in the “address” tag of in the xml input file and in the input of the client.

Unfortunately, UNIX sockets do not allow one to run i-PI and the clients on different computers, which limits greatly their utility when one needs to run massively parallel calculations. In these cases – typically when performing *ab initio* simulations – the force calculation becomes the bottleneck, so there is no need for fast communication with the server, and one can use internet sockets, that instead are specifically designed for communication over a network.

Internet sockets are described by an address and a port number. The address of the host is given as the IP address, or as a hostname that is resolved to an IP address by a domain name server, and is specified by the “address” variable of a object. The port number is an integer between 1 and 65535 used to distinguish between all the different sockets open on a particular host. As many of the lower numbers are protected for use in important system processes or internet communication, it is generally advisable to only use numbers in the range 1025-65535 for simulations.

The object has two more parameters. The option “latency” specifies how often i-PI polls the list of active clients to dispatch positions and collect results: setting it to a small value makes the program more responsive, which is appropriate when the evaluation of the forces is very fast. In *ab initio* simulations, it is best to set it to a larger value (of the order of 0.01 seconds), as higher latency will have no noticeable impact on performance, but will reduce the cost of having i-PI run in the background to basically zero.

Normally, i-PI can detect when one of the clients dies or disconnects, and can remove it from the active list and dispatch its force calculation to another instance. If however one of the client hangs without closing the communication channel, i-PI has no way of determining that something is going wrong, and will just wait forever. One can specify a parameter “timeout”, that corresponds to the maximum time – in seconds – that i-PI should wait before deciding that one of the clients has become unresponsive and should be discarded.

Running i-PI over the network

Understanding the network layout

Running i-PI in any non-local configuration requires a basic understanding of the layout of the network one is dealing with. Each workstation, or node of a HPC system, may expose more than one network interface, some of which can be connected to the outside internet, and some of which may be only part of a local network. A list of the network interfaces available on a given host can be obtained for instance with the command

```
> /sbin/ip addr
```

which will return a list of interfaces of the form

Each item corresponds to a network interface, identified by a number and a name (lo, eth0, eth1, ...). Most of the interfaces will have an associated IP address – the four numbers separated by dots that are listed after “inet”, e.g. 192.168.1.254 for the eth0 interface in the example above.




Fig. 5: A schematic representation of the network layout one typically finds when running i-PI and the clients on a HPC system and/or on a local workstation.

Figure 3.1 represents schematically a typical network layout for a HPC system and a local workstation. When running i-PI locally on a workstation, one can use the loopback interface (that can be referred to as “localhost” in the “address” field of both i-PI and the client) for communication. When running both i-PI and the clients on a HPC cluster, one should work out which of the the interfaces that are available on the node where the i-PI server runs are accessible from the compute nodes. This requires some trial and error, and possibly setting the “address” field dynamically from the job that launches i-PI. For instance, if one was running i-PI on the login node, and the clients on different compute nodes, as in Figure 2.1b, then on the HPC system described in Figure 3.1 one should set the address to that of the *ib1* interface – 111.111.111.111 in the example above. If instead i-PI was launched in a job script, then the submission script would have to check for the IP address associated with the *ib0* interface on the node the job has been dispatched to, and set that address (e.g. 111.111.111.200) in the inputs of both i-PI and the clients that will be launched in the same (or separate) jobs.

Running i-PI on a separate workstation (Figure 2.1c) gives maximum flexibility, but is also trickier as one has to reach the internet from the compute nodes, that are typically not directly connected to it. We discuss this more advanced setup in the next paragraph.

ssh tunnelling

If i-PI is to be run in a distributed computing mode, then one should make sure that the workstation on which the server will run is accessible from the outside internet on the range of ports that one wants to use for i-PI. There are ways to circumvent a firewall, but we will not discuss them here, as the whole point of i-PI is that it can be run on a low-profile PC whose security does not need to be critical. Typically arrangements can be made to open up a range of ports for incoming connections.

A more substantial problem – as it depends on the physical layout of the network rather than on software settings of the firewall – is how to access the workstation from the compute nodes, which in most cases do not have a network interface directly connected to the outside internet.

The problem can be solved by creating a ssh tunnel, i.e. an instance of the ssh secure shell that will connect the compute node to the login node, and then forward all traffic that is directed to a designated port on the compute node to the remote location that is running i-PI, passing through the outbound network interface of the login node.

In the example above, if i-PI is running on a local workstation, one should run:

from the job script that launches the client. For instance, with the network layout of Figure 3.1, and if the i-PI server is listening on port 12345 of the *eth0* interface, the tunnel should be created as:

```
> ssh -f -N -L 54321:123.123.123.123:12345 -2 111.111.111.111
```

The client should then be configured to connect to *localhost* on port 54321. The connection with i-PI will be established through the tunnel, and the data exchange can begin.

Note that, in order to be able to include the above commands in a script, the login node and the compute nodes should be configured to allow password-less login within the HPC system. This can be achieved easily, and does not entail significant security risks, since it only allows one to connect from one node to another within the local network. To do so, you should log onto the HPC system, and create a pair of ssh keys (if this has not been done already, in which case an *id_rsa.pub* file should be present in the user's *~/.ssh/* directory) by issuing the command

```
> ssh-keygen -t rsa
```

The program will then prompt for a passphrase twice. Since we wish to have use this in a job script where we will not be able to enter a password, just hit enter twice.

This should now have created two files in the directory *~/.ssh/*, *id_rsa* and *id_rsa.pub*. These should be readable only by you, so use the following code to set up the correct file permissions:

Finally, copy the contents of the file *id_rsa.pub* and append them to the file *authorized_keys* in the directory *~/.ssh/* of the user on the login node, which is typically shared among all the nodes of a cluster and therefore allows password-less login from all of the compute nodes.

4 A simple tutorial

Here we give a simple step-by-step guide through an example simulation, exploring some of the more generally useful options that i-PI offers and making no assumptions of previous experience of this code or other MD codes. Excerpts from the relevant input files are reproduced here, for explanation purposes, but to get the most out of this tutorial the user is strongly encouraged to work through it themselves. For this purpose, the input files have been included with the i-PI distribution, in the “test/tutorial” directory.

The chosen problem is that of a small *NPT* simulation of para-hydrogen, using the Silvera-Goldman potential [SG78]. We will take $(N, P, T) = (108, 0, 25 \text{ K})$.

4.1 Part 1 - *NVT* Equilibration run

Client code

Let us now consider the problem of how to use i-PI to run a *NPT* simulation of para-hydrogen. The first thing that is required is a client code that is capable of calculating the potential interactions of para-hydrogen molecules. Fortunately, one of the client codes distributed with i-PI has an appropriate empirical potential already hard-coded into it, and so all that is required is to create the “i-pi-driver” file compiling the code in the “drivers” directory, using the UNIX utility *make*.

This client code can be used for several different problems (see 2.3.2.1), some of which are explored in the “examples” directory, but for the current problem we will use the Silvera-Goldman potential with a cut-off radius of $15 a_0$. This is run using the following command:

```
> i-pi-driver -m sg -h localhost -o 15 -p 31415
```


The option “-m” is followed by the empirical potential required, in this case we use “sg” for Silvera-Goldman, “-h localhost” sets up the client hostname as “localhost”, “-o 15” sets the cut-off to 15 a_0 , and “-p 31415” sets the port number to 31415.

Note that usually this step will require setting up appropriate client code input files, possibly for an *ab initio* electronic structure code, and so is generally a more involved process. Refer to 2.3.2, and the documentation of the appropriate client code, for more details on how to do this step.

Creating the xml input file

Now that the client code is ready, an appropriate xml input file needs to be created from which the host server and the simulation data can be initialized. Here, we will go step by step through the creation of a minimal input file for a simple *NVT* equilibration run. Note that the working final version is held within the “tutorial-1” directory for reference.

Firstly, when reading the input file the i-PI xml functions look for a “simulation” tag as a sign to start reading data. For those familiar with xml jargon, we have defined as the root tag, so all the input data read in must start and end with a tag, as show below:

```
<simulation> Input data here... </simulation>
```

xml syntax requires a set of hierarchially nested tags, each of which contain data and/or more tags. Also, i-PI itself requires certain tags to be present, and keeps track of which tags are supposed to be where. More information about which tags are available can be found in 4, more information on xml syntax can be found in 3.1.1, and possible errors which can occur if the input file is not well formed can be found in 6.

For the sake of this first tutorial however, we will simply discuss the those tags which are needed for a single *NVT* equilibration run. The most important tags are , , , “total_steps” , and . These correspond to the tag to initialize the atom configurations (initialize), the tag ensemble defines the properties of the statistical ensemble considered (like temperature and pressure). While motion contains everything regarding the motion of the atoms in general and allows to specify which type of computation is expected (dynamics, geometry optimization, etc.), the tag dynamics contains all the specifics relative to the molecular dynamics (such as the barostat, thermostat and the time step, whose total number is defined in the total_step tags.). In theory it would be possible to join many force with different way into the tag forces, even though each of those forces needs a force field socket (ffsocket tag) that specify who is in charge of computing the forces and sending them to i-PI. Note that the attribute “forcefield” of the force tag must correspond to the attribute “name” of one of the ffsocket tag. All the tags that define the actual simulation must be within the environment. We will also discuss , which is used to define what output data is generated by the code.

At this point then, the input file looks like:

Initializing the configurations

Now let us consider each of these tags in turn. Firstly, . As the name suggests, this initializes the state of the system, so this is where we will specify the atom positions and the cell parameters. Firstly, this takes an attribute which specifies the number of replicas of the system, called “nbeads”. An attribute is a particular type of xml syntax designed to specify a single bit of data, and has the following syntax:

```
<initialize nbeads='4'> ... </initialize>
```

Note that an attribute forms part of the opening tag, and that the value being assigned to it is held within quotation marks. In this case, we have set the number of replicas, or beads, to 4.

Next, we must specify the atomic configuration. Rather than initialize the atom positions manually, we will instead use a separate configuration file for this purpose. Here we will discuss two of the input formats that are compatible with i-PI: xyz files and pdb files.

Note that, for the sake of this tutorial, we have included valid xyz and pdb input files in the “tutorial-1” directory called “our_ref.xyz” and “our_ref.pdb” respectively.

The xyz format is the simplest input format for a configuration file that i-PI accepts, and has the following syntax:

```
natoms # COMMENT LINE + CELLabcABC: a b c A B C cellangstrom
positionsangstrom atom1 x1 y1 z1 atom2 x2 y2 z2 ...
```

where “natoms” is replaced by an integer giving the total number of atoms, in this case 108, atom1 is a label for atom 1, in this case H2 (since we are simulating para-hydrogen), and (x1, y1, z1) are the x, y and z components of atom 1 respectively. The comment line can also contain the cell parameters and the position and cell units (angstrom in the example above).

Note that we are treating the para-hydrogen molecules isotropically here, i.e. as spherical pseudo-atoms. For the current system this is a good approximation, since at the state point under consideration every molecule is in its rotational ground state. For further details on this potential, and a demonstration of its application to quantum dynamics, see [SG78] and [MM05].

Other than its simplicity, the main advantage of this type of file is that it is free-formatted, and so there is no set precision to which each value must be written. This greatly simplifies both reading and writing these files.

The other file format that we can use is the pdb format. This has the following structure:

```
TITLE insert title here... positionangstrom cellangstrom CRYST1 a b c
A B C P 1 1 ATOM 1 n1 1 1 x1 y1 z1 0.00 0.00 0 ATOM 2 n2 1 1 x2 y2 z2
0.00 0.00 0 ...
```

where a, b and c are the cell vector lengths, A, B and C are the angles between them, n1 and n2 are the labels for atoms 1 and 2, and (x1, y1, z1) and (x2, y2, z2) give the position vectors of atoms 1 and 2.

Note that this is fixed-formatted, so the number of spaces matters. Essentially, the above format needs to be copied verbatim, using the same column widths and all the same keywords. For an exact specification of the file format (of which only a subset is implemented with i-PI) see http://deposit.rcsb.org/adit/docs/pdb_atom_format.html.

Here we will show how to specify the xml input file in both of these cases, assuming that the user has already created the configuration file themselves. Note that these file formats can be read by visualization programs such as VMD, and so it is generally advised when making your own input files to use such software to make sure that the configuration is as expected.

To use a configuration file the tag `in` should be used. This will take an input file with a given name and use it to initialize all relevant data. Both of these formats have the atom positions and labels, so this will initialize the positions, labels and masses of all the particles in the system, with the masses being implicitly set based on the atom label. The pdb configuration file will also be used to set the cell parameters.

Let us take these two file types in turn, and form the appropriate input sections. First, the xyz file. There are two attributes which are relevant to the tag for our current problem, “mode” and “units”. “mode” is used to describe what kind of data is being used to initialize from, and so in this case will be “xyz”. “units” specifies which units the file is given in, and so in this case is given by “angstrom”, which are the standard units of both xyz and pdb files. Note that if no units are specified then atomic units are assumed. For more information on the i-PI unit conversion libraries, and the available units, see 3.1.1.1.

The “units” attribute is now deprecated and will be removed in the future version of i-PI. The alternative, and the only one available in the future, is to specify the units within the comment line of the xyz or the TITLE line of the pdb formats (as shown in the examples above). It is also important to put the units only in one place: if the units will be present in both, the configuration file with the tag “units” and in the input files (xyz or pdb) the conversion will be applied twice. It is also important to note that the units of the cell parameters and the units of the content of the files are specified separately (“positionunits” specify the units of the data and “cellunits” specify the units of the cell). This is necessary because the xyz format can be used to store also quantity which have a different dimension than length

(velocities, forces, etc.). Even the cell parameters can now be specified directly within the xyz format. The comment line is parsed looking for a cell specification in of the following format:

- “CELL{abcABC}::” followed by six float numbers.
- “CELL{H}::” followed by nine float numbers.
- “CELL{GENH}::” followed by nine float numbers.

The “CELL{abcABC}” must be followed by the length of the vector cell and the three angle between them (as in the CRYST1 field of the pdb format -see above-). The other two must be followed by nine floats specifying, respectively, all the values of the cell matrix (flattened) or all the value of the inverse of the cell matrix (flattened).

Since the units are already specified into the xyz and pdb files, the config file will contain:

If the cell parameters are not specified in the xyz file, then, in the configuration file we must specify them separately. To initialize just the cell parameters, we use the tag . These could in theory be set using a separate file, but here we will initialize them manually. Taking a cubic cell with cell parameter 17.847 angstroms, we can specify this using the tag in three different ways:

```
<cell mode='manual' units='angstrom'> [17.847, 0, 0, 0, 17.847, 0, 0, 0, 17.847] </cell>
```

```
<cell mode='abcABC' units='angstrom'> [17.847, 17.847, 17.847, 90, 90, 90] </cell>
```

```
<cell mode='abc' units='angstrom'> [17.847, 17.847, 17.847] </cell>
```

If the xyz already contains the cell parameters, i-PI will use those which are read the last in the config file (if the “cell” tag follows the “file” specification then the cell parameters are those defined in the “cell” tag. If, otherwise, the “cell” tag compares in the config file before the “file” specification, then the cell parameters of the xyz file are used).

Note the use of the different “mode” attributes, “manual”, “abcABC” and “abc”. The first creates the cell vector matrix manually, the second takes the length of the three unit vectors and the angles between them in degrees, and the last assumes an orthorhombic cell and so only takes the length of the three unit vectors as arguments. We will take the last version for brevity, giving as our final section:

The pdb file is specified in a similar way, except that no tag needs be specified and the “mode” tag should be set to “pdb” (the units should be specified into the pdb file as shown in the example above):

As well as initializing all the atom positions, this section can also be used to set the atom velocities. Rather than setting these manually, it is usually simpler to sample these randomly from a Maxwell-Boltzmann distribution. This can be done using the tag by setting the “mode” attribute to “thermal”. This then takes an argument specifying the temperature to initialize the velocities to. With this, the final section is:

Creating the server socket

Next let us consider the and the sections, which deals with communication with the client codes. Since in this example we do not use ring-polymer contraction, we only need to specify a single tag:

```
<ffsocket> ... </ffsocket>
```

A socket is specified with three parameters; the port number, the hostname and whether it is a unix or an internet socket. These are specified by the “port” and “address” tags and the “mode” attribute respectively. To match up with the client socket specified above, we will take an internet socket on the hostname localhost and use port number 31415.

This gives the final section:

```
<ffsocket mode="inet" name="driver-sg"> <address> localhost
</address> <port> 31415 </port> </ffsocket>
```

To enhance the generality of the input, the forces used to move the atoms are read from the “forces” tag, within the “system” environment:

```
<system> ... <forces> ... </forces> ... </system>
```

Within this tag, the user can specify many different “force” tags. The final force will be the sum of the contribution from each “force” tag. Each force tag must be associated to an “ffsocket”. In particular the “forcefield” attribute of the “force” must match the “name” attribute of the “ffsocket”. In this tutorial only a single force is used that must match the ffsocket created above:

```
<system> ... <forces> <force forcefield="driver-sg"/> </forces> ...
</system>
```

For the sake of this tutorial, the “force” tag must be empty.

Generating the correct dynamics

The next section that we will need is , which determines the computation i-pi will perform. Since we wish to do a molecular dynamics, the attribute “mode” of the “motion” tag must be equal to “dynamics”. The details of the dynamics integration are given within . Since we wish to do a *NVT* simulation, we set the “mode” attribute to “nvt” (note that we use lower case, and that the tags are case sensitive), and must specify the temperature using the appropriate tag:

```
<motion mode='dynamics'> <dynamics mode='nvt'> ... </dynamics>
</motion>
```

This defines the computation that will be performed. We also must decide which integration algorithm to use, and how large the time step should be. In general, the time step should be made as large as possible without there being a drift in the conserved quantity. Usually we would take a few short runs with different time steps to try and optimize this, but for the sake of this tutorial we will use a safe value of 1 femtosecond, giving:

```
<dynamics mode='nvt'> ... <timestep units='femtosecond'> 1
</timestep> </dynamics>
```

Finally, while the microcanonical part of the integrator is initialized automatically, there are several different options for the constant temperature sampling algorithm, specified by . For simplicity we will use (the global version of) the path-integral Langevin equation (PILE) algorithm [CPMM10], which is specifically designed for path integral simulations. This is specified by the “mode” tag “pile_g”. This integrator also has to be initialized with a time scale parameter, “tau”, which determines how strong the thermostat is, which we will set to 25 femtoseconds. Putting all of this together, we get:

Now that we have decided on the time step, we will decide the total number of steps to run the simulation for. Equilibrating the system is likely to take around 5 picoseconds, so we will take 5000 time steps, using:

The temperature must be specified within the :

```
<system> ... <ensemble> <temperature units='kelvin'> 300
</temperature> </ensemble> ... </system>
```

Customizing the output

So far, we have only considered how to set up the simulation, and not the data we wish to obtain from it. However, there are a wide variety of properties of interest that i-PI can calculate and a large number of different output options, so to avoid confusion let us go through them one at a time.

Firstly, the amount of data sent to standard output can be adjusted with the “verbosity” attribute of :

```
<simulation verbosity='high'> ... </simulation>
```

By default the verbosity is set to “low”, which only outputs important warning messages and information, and some statistical information every 1000 time steps. Here we will set it to “high”, which will tell i-PI to output the following data every time step:

```
# Average timings at MD step S. t/step: TOTAL [p: P q: Q t: T] # MD
diagnostics: V: POTENTIAL Kcv: KINETIC Ecns: CONSERVED @SOCKET:
Assigning [ X ] request id ID to client with last-id LID ( CID/ CTOT
: )
```

where the output values have been replaced with the following:

S:

This gives the current time step.

TOTAL:

This gives the amount of time the current time step took.

P:

This gives how long the momentum propagation step took.

Q:

This gives how long the free-ring polymer propagation step took.

T:

This gives how long the thermostat integration step took.

POTENTIAL:

This gives the current potential energy of the system.

KINETIC:

This gives the current kinetic energy of the system.

CONSERVED:

This gives the current conserved quantity.

X:

This says whether or not i-PI found a match for the calculation of replica ID or not. If one of the connected client codes calculated the forces for this replica on the last time step, then X will be “match”, and i-PI will automatically assign this replica to the same client as before. This should happen with all the replicas if CTOT is the same as the number of beads.

ID:

The index of the replica currently being assigned to a client code.

LID:

The index of the replica which the client code last did a force calculation of.

CID:

The index of the client code in the list of all connected client codes.

CTOT:

The total number of connected client codes.

What output gets written to file is specified by the tag. There are three types of files; properties files, trajectory files and checkpoint files, which are specified with , and tags respectively. For an in-depth discussion on these three types of output files see 3.2, but for now let us just explain the rationale behind each of these output file types in turn.

checkpoint files:

These give a snapshot of the state of the simulation. If used as an input file for a new i-PI simulation, this simulation will start from the point where the checkpoint file was created in the old simulation.

trajectory files:

These are used to print out properties relevant to all the atoms, such as the velocities or forces, for each degree of freedom. These can be useful for calculating correlation functions or radial distribution functions, but possibly their most useful feature is that visualization programs such as VMD can read them, and then use this data to show a movie of how the simulation is progressing.

properties files:

These are usually used to print out system level properties, such as the timestep, temperature, or kinetic energy. Essentially these are used to keep track of a small number of important properties, either to visualize the progress of the simulation using plotting programs such as gnuplot, or to be used to get ensemble averages.

Now that we know what each input file is used for, let us take an example of an output section and show how the xml input section works. The default output, i.e. what would be output if nothing was set by the user, would be generated with the following section:

This creates 6 files: “i-pi.md”, “i-pi.pos_0.xyz”, “i-pi.pos_1.xyz”, “i-pi.pos_2.xyz”, “i-pi.pos_3.xyz” and “i-pi.checkpoint”. “i-pi.md” is the properties file, “i-pi.pos_x.xyz” are the position trajectory files, and “i-pi.checkpoint” is the checkpoint file.

The filenames are created using the syntax “prefix”.“filename”[_(file specifier)][.(file format)], where the file specifier is added to separate similar files. For example, in the above case the different position trajectories for each bead are given a file specifier corresponding to the appropriate bead index.

The “stride” attributes set how often data is output to each file; so in the above case the properties are written out every 10 time steps, the trajectories every 100, and the checkpoints every 1000. The “format” attribute sets the format of the trajectory files, and the “overwrite” attribute sets whether each checkpoint file overwrites the previous one or not.

There are several options we can use to customize the output data. Firstly, the “prefix” attribute should be set to something which can be used to distinguish the files from different simulation runs. In this case we can simply set it to “tut1”:

```
<output prefix='tut1'> ... </output>
```

As for the input parameters, the units the output data is given in can be set by the user. Unlike the input parameters however, this is done by specifying an appropriate unit in curly braces after the name of the property or trajectory of interest, as shown below:

Next, let us adjust some of the attributes. Let us suppose that we wish to output the properties every time step, to check for conserved quantity jumps, and to output the trajectory in pdb format. To do this we would set the “stride” and “format” tags, as shown below:

Note that we have added a “cell_units” attribute to the tag, so that the cell parameters are consistent with the position output.

Finally, let us suppose that we wished to output another output property to a different file to the others. One example of when this might be necessary is if there were an output property which was more expensive to calculate than the others, and so it would be impractical to output it every time step. With i-PI this is easy to do, all that is required is to add another tag with a different filename.

For demonstration purposes, we will choose to print out the forces acting on one tagged bead, since this requires an argument to be passed to the function that calculates it. The i-PI syntax for doing this is to have the arguments to be passed to the function between standard braces, separated by semi-colons.

To print out the forces acting on one bead we need the “atom_f” property, which takes two arguments, “atom” and “bead”, giving the index of the atom and bead tagged respectively. The appropriate syntax is then given below:

This will print out the force vector acting on bead 0 of atom 0. i-PI also accepts positional arguments (i.e. arguments not specified by a name, but just by their position in the list of arguments), and so this could also be written as:

Finally, putting all this together, and adjusting some of the parameters of the new file, we get:

Running the simulation

Now that we have a valid input file, we can run the test simulation. The “i-pi” script in the root directory is used to create an i-PI simulation from a xml input file. As explained in 2.3.1 this script is run (if we assume that we are in the “tutorial-1” directory) using:

```
> python ../../../../i-pi tutorial-1.xml
```

This will start the i-PI simulation, creating the server socket and initializing the simulation data. This should at this point print out a header message to standard output, followed by a few information messages that end with “starting the polling thread main loop”, which signifies that the server socket has been opened and is waiting for connections from client codes.

At this point the driver code is run in a new terminal from the “drivers” directory using the command specified above:

```
> i-pi-driver -m sg -h localhost -o 15 -p 31415
```

The i-PI code should now output a message saying that a new client code has connected, and start running the simulation.

Output data

Once the simulation is finished (which should take about half an hour) it should have output “tut1.md”, “tut1.force”, “tut1.pos_0.xyz”, “tut1.pos_1.xyz”, “tut1.pos_2.xyz”, “tut1.pos_3.xyz”, “tut1.checkpoint” and “RESTART”.

Firstly, we consider the checkpoint files, “tut1.checkpoint” and “RESTART”. As mentioned before, these files can be used as a means of restarting the simulation from a previous point. As an example, the last checkpoint should have been at step 4999, and so we could rerun the last step using

```
> ../../../../i-pi tut1.checkpoint
```

followed by running “i-pi-driver” as before.

The difference between these two files is that, while “tut1.checkpoint” was specified by the user, “RESTART” is automatically generated at the end of every i-PI run. This file then is what we will need to initialize the *NPT* run, since it contains the state of the system after equilibration.

Next, let us look at the trajectory files. Since we have printed out the positions, these should tell us how the spatial distribution has equilibrated, and give us some insight into the atom dynamics. The easiest way to use these files, as discussed earlier, is to use the trajectory files as input to a visualization program such as VMD.

If we do this with these files, we see that the simulation started from a crystalline configuration and then over the course of the simulation began to melt. Since the state point studied and with the potential given para-hydrogen is a liquid [SG78], this is what we would expect.

Finally, let us check the “tut1.md” file. For the current problem, i.e. checking that we have a suitably equilibrated system, we should do two tests. Firstly, we should check that the conserved quantity does not exhibit any major drift,

and second we should check to see if the properties of interest have converged. Using gnuplot, we can plot the relevant graphs using:

```
> gnuplot > p './tut1.md' u 1:3 # Plots column 1, i.e. current
simulation step, > p './tut1.md' u 1:4 # against columns 3, 4, 5 and
6, > p './tut1.md' u 1:5 # i.e. conserved quantity, temperature, > p
'./tut1.md' u 1:6 # potential energy and kinetic energy
```

This will show that the conserved quantity has only a small drift upwards, the kinetic and potential energies have equilibrated, and the thermostat is keeping the temperature at the specified value. We have therefore specified a sufficiently short time step, chosen the thermostat parameters sensibly, and have equilibrated the properties of interest. Therefore this stage of the simulation is done, and we are ready to start the *NPT* run.

4.2 Part 2 - *NPT* simulation

Now that we have converged *NVT* simulation data, we can use this to initialize a *NPT* simulation. There are two ways of doing this, both of which involve using the RESTART file generated at the end of the *NVT* run as a starting point. Note that for simplicity we will again take $N = 108$, $T = 25K$, and use $P = 0$.

Modifying the RESTART file

Firstly, you can use the RESTART file directly, modifying it so that instead of continuing with the original *NVT* simulation it will instead start a new *NPT* simulation. We have included in the “tutorial-2” directory both a RESTART file from tutorial 1 and an adjusted file which will run *NPT* dynamics, “tutorial-2a.xml”

These adjustments start with resetting the “step” tag, so that it starts with the value 0. This can be done by simply removing the tag. Similarly, we can increase the total number of steps so that it is more suitable for collecting the necessary amount of *NPT* data, in this case we will set “total_steps” to 100000.

We will also update the output files, first by setting the filenames to start with “tut2a” rather than “tut1”, and secondly by adding the volume and pressure to the list of computed properties so that we can check that the ensemble is being sampled correctly. Putting this together this gives:

Finally, we must change the and the tags so that the correct ensemble is sampled. The first thing that must be done is adding a “pressure” tag in the ensemble:

```
<ensemble> <pressure> 0 </pressure> ... </ensemble>
```

Then, we must also specify the constant pressure algorithm, using the tag within the dynamics environment. Do not forget to change the mode attribute of the dynamics from “nvt” to “npt”. This example uses a stochastic barostat to apply pressure to an isotropic system, which can be specified with the option “isotropic”. See the documentation of the object and the examples to see how to apply an anisotropic stress, or to allow for cell shape fluctuations.

The isotropic barostat also requires a thermostat to deal with the volume degree of freedom, which we will take to be a simple Langevin thermostat. This thermostat is specified in the same way as the one which does the constant temperature algorithms for the atomic degrees of freedom, and we will take its time scale to be 250 femtoseconds:

```
<system> <ensemble> <pressure> 0 </pressure> </ensemble> <motion
mode='dynamics'> <dynamics mode='npt'> <barostat mode='isotropic'>
<thermostat mode='langevin'> <tau units='femtosecond'> 250 </tau>
</thermostat> ... </barostat> ... </dynamics> ... </motion> ...
</system>
```

Finally, we will take the barostat time scale to be 250 femtoseconds also, giving:


```
<system> <ensemble> <pressure> 0 </pressure> </ensemble> <motion
mode='dynamics'> <dynamics mode='npt'> <barostat mode='isotropic'>
<thermostat mode='langevin'> <tau units='femtosecond'> 250 </tau>
</thermostat> <tau units='femtosecond'> 250 </tau> </barostat> ...
</dynamics> ... </motion> ... </system>
```

with the rest of the and tag being the same as before.

Initialization from RESTART

A different way of initializing the simulation is to use the RESTART file as a configuration file, in the same way that the xyz/pdb files were used previously.

Firstly, the original input file “tutorial-1.xml” needs to be modified so that it will do a *NPT* simulation instead of *NVT*. This involves modifying the “total_steps”, and tags as above. Next, we replace the tag section with:

Note that the “mode” attribute has been set to “chk” to specify that the file is a checkpoint file. This will then use the RESTART file to initialize the bead configurations and velocities and the cell parameters.

Again, there is a file in the “tutorial-2” directory for this purpose, “tutorial-2b.xml”.

Running the simulation

Whichever method is used to create the input file, the simulation is run in the same way as before, using either “tutorial-2a.xml” or “tutorial-2b.xml” as the input file. Note how the volume fluctuates with time, as it is no longer held constant in this ensemble.

4.3 Part 3 - A fully converged simulation

As a final example, we note that at this state point 16 replicas and at least 172 particles are actually required to provide converged results. As a last tutorial then, you should repeat tutorials 1 and 2 with this number of replicas and atoms.

The directory “tutorial-3” contains *NVT* and *NPT* input files which can be used to do a fully converged *NPT* simulation from scratch, except that they are missing some of the necessary input parameters.

If these are chosen correctly and the simulation is run properly the volume will be 31 cm³/mol and the total energy should be -48 K [MHT99].

With this number of beads and atoms, the force calculation is likely to take much longer than it did in either tutorial 1 or 2. To help speed this up, we will now discuss how to parallelize the calculation over the sockets, and how to speed up the data transfer.

Firstly, in this simple case where we are calculating an isotropic, pair-wise interaction, the data transfer time is likely to be a significant proportion of the total calculation time. To help speed this up, there is the option to use a unix domain socket rather than an internet socket. These are optimized for local communication between processes on a single computer, and so for the current problem they will be much faster than internet sockets.

To specify this, we simply set the “mode” attribute of the tag to “unix”:

```
<ffsocket mode='unix' name="driver-sg"> ... </ffsocket>
```

We then specify that the client code should connect to a unix socket using the -u flag:

```
> i-pi-driver -u -m sg -h localhost -o 15 -p 31415
```

Parallelizing the force calculation over the different replicas of the system is similarly easy, all that is required is to run the above command multiple times. For example, if we wish to run 4 client codes, we would use:

```
> for a in 1 2 3 4; do > i-pi-driver -u -m sg -h localhost -o 15 -p  
31415 & > done
```

Using these techniques should help speed up the calculation considerably, at least in this simple case. Note however, that using unix domain sockets would give a negligible gain in speed in most simulations, since the force calculation usually takes much longer than the data transfer.

5 Input file reference

This chapter gives a complete list of the tags that can be specified in the xml input file, along with the hierarchy of objects. Note that every xml input file must start with the root tag . See the accompanying “help.xml” file in the “doc/help_files” directory to see the recommended input file structure.

Each section of this chapter describes one of the major input classes used in the i-PI initialization. These sections will start with some normal text which describes the function of this class in detail. After this the attributes of the tag will be listed, enclosed within a frame for clarity. Finally, the fields contained within the tag will be listed in alphabetical order, possibly followed by their attributes. **Attribute** names will be bold and **field** names both bold and underlined. *Other supplementary information* will be in small font and italicized.

5.1 simulation

This is the top level class that deals with the running of the simulation, including holding the simulation specific properties such as the time step and outputting the data.

ATTRIBUTES

verbosity

The level of output on stdout.

dtype: string

options: ['quiet', 'low', 'medium', 'high', 'debug']

default: medium

threading

Whether multiple-systems execution should be parallel. Makes execution non-reproducible due to the random number generator being used from concurrent threads.

dtype: boolean

default: True

mode

What kind of simulation should be run.

dtype: string

options: ['md', 'static']

default: md

safe_stride

Consistent simulation states will be saved every this number of steps. Saving state entails a small overhead, so you may want to set this to the smallest output frequency in your simulation to make i-PI faster. Use at your own risk!

dtype: integer

default: 1

floatformat

A format for all printed floats.

dtype: string

default: %16.8e

FIELDS

prng

Deals with the pseudo-random number generator.

output

This class defines how properties, trajectories and checkpoints should be output during the simulation. May contain zero, one or many instances of properties, trajectory or checkpoint tags, each giving instructions on how one output file should be created and managed.

step

The current simulation time step.

dtype: integer

default: 0

total_steps

The total number of steps that will be done. If 'step' is equal to or greater than 'total_steps', then the simulation will finish.

dtype: integer

default: 1000

total_time

The maximum wall clock time (in seconds).

dtype: float

default: 0

smotion

Options for a 'super-motion' step between system replicas

system

This is the class which holds all the data which represents a single state of the system.

system_template

Generic input value

ffsocket

Deals with the assigning of force calculation jobs to different driver codes, and collecting the data, using a socket for the data communication.

fflj

Simple, internal LJ evaluator without cutoff, neighbour lists or minimal image convention. Expects standard LJ parameters, e.g. { eps: 0.1, sigma: 1.0 }.

ffdmd

Simple, internal DMD evaluator without neighbor lists, but with PBC. Expects coupling elements ($n*(n-1)/2$ of them), oscillating frequency and time step.

ffdebye

Harmonic energy calculator

ffplumed

Direct PLUMED interface

ffyaff

Uses a Yaff force field to compute the forces.

ffcommittee

Combines multiple forcefields to build a committee model, that can be used to compute uncertainty-quantified machine-learning models. Each forcefield can be any of the other FF objects, and each should be used with a client that generates a slightly different estimation of energy and forces. These are averaged, and the mean used as the actual forcefield. Statistics about the distribution are also returned as extras fields, and can be printed for further postprocessing. Also contains options to use it for uncertainty estimation and for active learning in a ML context, based on a committee model. Implements the approaches discussed in DOI: 10.1063/5.0036522.

ffsgdml

A SGDML energy calculator

ffcavphsocket

A cavity molecular dynamics driver for vibrational strong coupling. In the current implementation, only a single cavity mode polarized along the x and y directions is coupled to the molecules. Check <https://doi.org/10.1073/pnas.2009272117> and also examples/lammps/h2o-cavmd/ for details.

5.2 system

This is the class which holds all the data which represents a single state of the system.

ATTRIBUTES

prefix

Prepend this string to output files generated for this system.

dtype: string

default:

FIELDS

initialize

Specifies the number of beads, and how the system should be initialized.

forces

Deals with creating all the necessary forcefield objects.

ensemble

Holds all the information that is ensemble specific, such as the temperature and the external pressure.

motion

Allow choosing the type of calculation to be performed. Holds all the information that is calculation specific, such as geometry optimization parameters, etc.

beads

Describes the bead configurations in a path integral simulation.

normal_modes

Deals with the normal mode transformations, including the adjustment of bead masses to give the desired ring polymer normal mode frequencies if appropriate. Takes as arguments frequencies, of which different numbers must be specified and which are used to scale the normal mode frequencies in different ways depending on which 'mode' is specified.

cell

Deals with the cell parameters. Takes as array which can be used to initialize the cell vector matrix.

dtype: float

dimension: length

default:

[0. 0. 0. 0. 0. 0. 0. 0. 0.]

5.3 system_template

Generic input value

FIELDS

template

A string that will be read verbatim containing the model for a system to be generated

dtype: string

labels

A list of strings that should be substituted in the template to create multiple systems

dtype: string

instance

A list of strings that should the labels creating one system instance

dtype: string

5.4 motion

Allow choosing the type of calculation to be performed. Holds all the information that is calculation specific, such as geometry optimization parameters, etc.

ATTRIBUTES

mode

How atoms should be moved at each step in the simulation. ‘replay’ means that a simulation is replayed from trajectories provided to i-PI.

dtype: string

options: ['vibrations', 'minimize', 'replay', 'neb', 'string', 'dynamics', 'constrained_dynamics', 't_ramp', 'p_ramp', 'alchemy', 'atomswap', 'planetary', 'instanton', 'al-kmc', 'dummy', 'scp', 'normalmodes', 'multi']

FIELDS

fixcom

This describes whether the centre of mass of the particles is fixed.

dtype: boolean

default: True

fixatoms

Indices of the atoms that should be held fixed.

dtype: integer

default: []

optimizer

Option for geometry optimization

neb_optimizer

Option for NEB optimization

string_optimizer

Option for String minimal-energy path optimization

dynamics

Option for (path integral) molecular dynamics

constrained_dynamics

Option for constrained classical molecular dynamics

file

This describes the location to read a trajectory file from. Replay syntax allows using some POSIX wildcards in the filename of trajectory files. If symbols `?*` are found in a filename, the code expects to find exactly Nbeads files that match the provided pattern. Bead indices will be read from the files, and the files will be ordered ascendingly by their bead indices. Wildcarded files are expected to be in the folder where the simulation runs.

dtype: string

default:

vibrations

Option for phonon computation

normalmodes

Option for solving the vibrational Schroedinger's equations in normal mode coordinates.

scp

Option for self consistent phonons computation

alchemy

Option for alchemical exchanges

atomswap

Option for Monte Carlo atom swap

t_ramp

Option for temperature ramp

p_ramp

Option for pressure ramp

instanton

Option for Instanton optimization

al6xxx_kmc

Option for Al-6xxx KMC

planetary

Option for planetary model calculator

motion

A motion class that can be included as a member of a 'multi' integrator.

5.5 cell

Deals with the cell parameters. Takes as array which can be used to initialize the cell vector matrix.

dtype: float

dimension: length

ATTRIBUTES

units

The units the input data is given in.

dtype: string

default: automatic

shape

The shape of the array.

dtype: tuple

default: (0,)

mode

If 'mode' is 'manual', then the array is read in directly, then reshaped according to the 'shape' specified in a row-major manner. If 'mode' is 'file' then the array is read in from the file given.

dtype: string

options: ['manual', 'file']

default: manual

5.6 smotion

Allow choosing the type of smotion to be performed. Holds all the information that is calculation specific, such as replica exchange parameters, etc.

ATTRIBUTES

mode

Kind of smotion which should be performed.

dtype: string

options: ['dummy', 'remd', 'metad', 'dmd', 'multi']

FIELDS

remd

Option for REMD simulation

metad

Option for REMD simulation

dmd

Option for driven MD simulation

smotion

A smotion class that can be included as a member of a 'multi' Smotion.

5.7 remd

Replica Exchange

FIELDS

stride

Every how often to try exchanges (on average).

dtype: float

default: 1.0

krescale

Rescale kinetic energy upon exchanges.

dtype: boolean

default: True

swapfile

File to keep track of replica exchanges

dtype: string

default: remd_idx

repindex

List of current indices of the replicas compared to the starting indices

dtype: integer

default: []

5.8 metad

MetaDynamics

FIELDS

metaff

List of names of forcefields that should do metadynamics.

dtype: string

default: []

use_energy

Transfer the potential energy value to PLUMED to use as a collective variable. Can only be used with classical simulations because it requires a rather hacky mechanism to transfer the energy of the system to the forcefield.

dtype: boolean

default: False

5.9 dmd

DrivenMD with external time-dependent driving potential

FIELDS

dmdff

List of names of forcefields that should do driven MD. Accepts ffdmd forcefield types. Currently implemented (2021) for ffdmd only the driving potential similar to the one described in Bowman, .., Brown JCP 119, 646 (2003).

dtype: string

default: []

5.10 ensemble

Holds all the information that is ensemble specific, such as the temperature and the external pressure.

FIELDS

temperature

The temperature of the system.

dtype: float

dimension: temperature

default: -1.0

pressure

The external pressure.

dtype: float

dimension: pressure

default: -12345.0

stress

The external stress.

dtype: float

dimension: pressure

default:

[-12345. -0. -0. -0. -12345. -0. -0. -0. -12345.]

eens

The ensemble contribution to the conserved quantity.

dtype: float

dimension: energy

default: 0.0

bias

Deals with creating all the necessary forcefield objects.

bias_weights

Bias weights.

dtype: float

default: []

hamiltonian_weights

Hamiltonian weights.

dtype: float

default: []

time

The internal time for this system

dtype: float

dimension: time

default: 0.0

5.11 bias

Deals with creating all the necessary forcefield objects.

FIELDS

force

The class that deals with how each forcefield contributes to the overall potential, force and virial calculation.

5.12 dynamics

Holds all the information for the MD integrator, such as timestep, the thermostats and barostats that control it.

ATTRIBUTES

mode

The ensemble that will be sampled during the simulation.

dtype: string

options: ['nve', 'nvt', 'npt', 'nst', 'sc', 'scnpt']

default: nve

splitting

The Louiville splitting used for sampling the target ensemble.

dtype: string

options: ['obabo', 'baoab']

default: obabo

FIELDS

thermostat

The thermostat for the atoms, keeps the atom velocity distribution at the correct temperature.

barostat

Simulates an external pressure bath.

timestep

The time step.

dtype: float

dimension: time

default: 1.0

nmts

Number of iterations for each MTS level (including the outer loop, that should in most cases have just one iteration).

dtype: integer

default: []

5.13 constrained_dynamics

Holds all the information for the MD integrator, such as timestep, the thermostats and barostats that control it.

ATTRIBUTES

mode

The ensemble that will be sampled during the simulation.

dtype: string

options: ['nve', 'nvt']

default: nve

splitting

The integrator used for sampling the target ensemble.

dtype: string

options: ['obabo', 'baoab']

default: baoab

FIELDS

thermostat

The thermostat for the atoms, keeps the atom velocity distribution at the correct temperature.

barostat

Simulates an external pressure bath.

timestep

The time step.

dtype: float

dimension: time

default: 1.0

nmts

Number of iterations for each MTS level (including the outer loop, that should in most cases have just one iteration).

dtype: integer

default: []

nsteps_o

The number of sub steps used in the evolution of the thermostat (used in function `step_Oc`). Relevant only for GLE thermostats

dtype: integer

default: 1

nsteps_geo

The number of sub steps used in the evolution of the geodesic flow (used in function `step_Ag`).

dtype: integer

default: 1

csolver

Define a numerical method for computing the projection operators associated with the constraint.

constraint

Define a constraint to be applied onto atoms

5.14 csolver

Holds all parameters for the numerical method used to solve the constraint.

FIELDS

tolerance

Tolerance value used in the Quasi-Newton iteration scheme.

dtype: float

default: 0.0001

maxit

Maximum number of steps used in the Quasi-Newton iteration scheme.

dtype: integer

default: 1000

norm_order

Order of norm used to determine termination of the Quasi-newton iteration.

dtype: integer

default: 2

5.15 constraint

Generic input value

ATTRIBUTES

mode

The type of constraint.

dtype: string

options: ['distance', 'angle', 'eckart', 'multi', 'multi']

default: distance

FIELDS

atoms

List of atoms indices that are to be constrained.

dtype: integer

default: []

values

List of constraint lengths.

dtype: float

dimension: length

default: []

constraint

One or more constraints that have to be considered coupled

5.16 t_ramp

TemperatureRamp Motion class. It just updates the ensemble temperature in steps, between the indicated temperatures, and then holds to the highest value. It should typically be combined with a dynamics class and thermostats, using a MultiMotion.

FIELDS

t_start

Initial temperature

dtype: float

dimension: energy

default: 1.0

t_end

Final temperature

dtype: float

dimension: energy

default: 1.0

logscale

Change temperature on a logarithmic scale.

dtype: boolean

default: False

total_steps

Total number of steps for the ramp

dtype: integer

default: 0

current_step

Current step along the ramp

dtype: integer

default: 0

5.17 p_ramp

PressureRamp Motion class. It just updates the ensemble pressure in steps, between the indicated values, and then holds to the highest value. It should typically be combined with a dynamics class and barostats, using a MultiMotion.

FIELDS

p_start

Initial pressure

dtype: float

dimension: pressure

default: 1.0

p_end

Final pressure

dtype: float

dimension: pressure

default: 1.0

logscale

Change pressure on a logarithmic scale.

dtype: boolean

default: False

total_steps

Total number of steps for the ramp

dtype: integer

default: 0

current_step

Current step along the ramp

dtype: integer

default: 0

5.18 alchemy

Holds all the information for doing Monte Carlo alchemical exchange moves.

ATTRIBUTES

mode

dtype: string

options: ['dummy']

default: dummy

FIELDS

names

The names of the atoms to be exchanged, in the format [name1, name2, ...].

dtype: string

default: []

nxc

The average number of exchanges per step to be attempted

dtype: float

default: 1

ealc

The contribution to the conserved quantity for the alchemical exchanger

dtype: float

default: 0.0

5.19 planetary

Holds all the information for the planetary model frequency matrix calculator.

ATTRIBUTES

mode

The constrained-centroid sampling mode.

dtype: string

options: ['md']

default: md

FIELDS

thermostat

The thermostat for the atoms, keeps the atom velocity distribution at the correct temperature.

timestep

The time step.

dtype: float

dimension: time

default: 1.0

nmts

Number of iterations for each MTS level (including the outer loop, that should in most cases have just one iteration).

dtype: integer

default: []

nsamples

Number of samples to accumulate for each planetary step.

dtype: integer

default: 0

stride

How often the planetary calculation should actually be triggered.

dtype: integer

default: 1

nbeads

Number of beads for centroid-constrained dynamics (default same as master trajectory)

dtype: integer

default: -1

screen

Screening parameter for path-integral frequency matrix.

dtype: float

dimension: length

default: 0.0

5.20 atomswap

Holds all the information for doing Monte Carlo atom swap moves.

ATTRIBUTES

mode

Dummy attribute, does nothing.

dtype: string

options: ['dummy']

default: dummy

FIELDS

names

The names of the atoms to be to exchanged, in the format [name1, name2, ...].

dtype: string

default: []

nxc

The average number of exchanges per step to be attempted

dtype: float

default: 1

ealc

The contribution to the conserved quantity for the atom swapper

dtype: float

default: 0.0

5.21 instanton

A class for instanton calculations

ATTRIBUTES

mode

Defines whether it is an instanton rate or instanton tunneling splitting calculaion

dtype: string

options: ['rate', 'splitting']

default: rate

FIELDS

tolerances

Convergence criteria for optimization.

biggest_step

The maximum step size during the optimization.

dtype: float

default: 0.4

old_pos

The previous step positions during the optimization.

dtype: float

dimension: length

default: []

old_pot

The previous step potential energy during the optimization

dtype: float

dimension: energy

default: []

old_force

The previous step force during the optimization

dtype: float

dimension: force

default: []

opt

The geometry optimization algorithm to be used. For small system sizes nichols is recommended. Lanczos is tailored for big bigger than nbonds*atoms >~38*64. NR works in both cases given that the initial guess is close to the optimized geometry. Finally lbfgs is used for tunneling splitting calculations.

dtype: string

options: ['nichols', 'NR', 'lbfgs', 'lanczos', 'None']

default: None

max_e

Evaluate the forces in a reduced ring polymer such that the potential energy between consecutive replicas is smaller than the provided value.

dtype: float

dimension: energy

default: 0.0

max_ms

Evaluate the forces in a reduced ring polymer such that that mass-scaled distance in a.u. between consecutive replicas is smaller than the provided value.

dtype: float

default: 0.0

discretization

Allows to specified non uniform time discretization as proposed in J. Chem. Phys. 134, 184107 (2011)

dtype: float

default: []

alt_out

Alternative output: Prints different formatting of outputs for geometry, hessian and bead potential energies. All quantities are also accessible from typical i-pi output infrastructure. Default to 1, which prints every step. -1 will suppress the output (except the last one). Any other positive number will set the frequency (in steps) with which the quantities are written to file. The instanton geometry is printed in xyz format and the distances are in angstroms The hessian is printed in one line with the following format: h1_1,h2_1,...,hN_1, h2_2,h2_2,hN_2, ,h1_d,h2_d,...,hN_d. Where N represents the total number of replicas, d the number of dimension of each replica (3*n_atoms) and hi_j means the row j of the physical hessian corresponding to the replica i. The physical hessian uses a convention according to the positions convention used in i-pi. Example of 2 particles, the first two rows of the physical hessian reads: 'H_x1_x1, H_x1_y1, H_x1_z1, H_x1_x2, H_x1_y2, H_x1_z2' 'H_x2_x1, H_x2_y1, H_x2_z1, H_x2_x2, H_x2_y2, H_x2_z2'

dtype: integer

default: 1

prefix

Prefix of the output files.

dtype: string

default: instanton

delta

Initial stretch amplitude.

dtype: float

default: 0.1

hessian_init

How to initialize the hessian if it is not fully provided.

dtype: string

options: ['true', 'false']

default: false

hessian

(Approximate) Hessian.

dtype: float

default: []

hessian_update

How to update the hessian after each step.

dtype: string

options: ['powell', 'recompute']

default: powell

hessian_asr

Removes the zero frequency vibrational modes depending on the symmetry of the system.

dtype: string

options: ['none', 'poly', 'crystal']

default: none

qlist_lbfgs

List of previous position differences for L-BFGS, if known.

dtype: float

default: []

glist_lbfgs

List of previous gradient differences for L-BFGS, if known.

dtype: float

default: []

old_direction

The previous direction in a CG or SD optimization.

dtype: float

default: []

scale_lbfgs

Scale choice for the initial hessian. 0 identity. 1 Use first member of position/gradient list. 2 Use last member of position/gradient list.

dtype: integer

default: 2

corrections_lbfgs

The number of past vectors to store for L-BFGS.

dtype: integer

default: 20

ls_options

Options for line search methods. Includes: tolerance: stopping tolerance for the search, iter: the maximum number of iterations, step: initial step for bracketing, adaptive: whether to update initial step.

energy_shift

Set the zero of energy.

dtype: float

dimension: energy

default: 0.0

hessian_final

Decide if we are going to compute the final big-hessian by finite difference.

dtype: string

options: ['false', 'true']

default: false

5.22 vibrations

Fill in.

ATTRIBUTES

mode

The algorithm to be used: finite differences (fd), normal modes finite differences (nmfd), and energy-scaled normal mode finite differences (enmfd).

dtype: string

options: ['fd', 'nmfd', 'enmfd']

default: fd

FIELDS

pos_shift

The finite displacement in position used to compute derivative of force.

dtype: float

default: 0.01

energy_shift

The finite displacement in energy used to compute derivative of force.

dtype: float

default: 0.0

output_shift

Shift by the dynamical matrix diagonally before outputting.

dtype: float

default: 0.0

prefix

Prefix of the output files.

dtype: string

default: phonons

asr

Removes the zero frequency vibrational modes depending on the symmetry of the system.

dtype: string

options: ['none', 'poly', 'lin', 'crystal']

default: none

dynmat

Portion of the dynamical matrix known up to now.

dtype: float

default: []

refdynmat

Portion of the refined dynamical matrix known up to now.

dtype: float

default: []

5.23 scp

Fill in.

ATTRIBUTES

mode

The statistics to be used in the calculation of the free energy. Quantum (qn) or classical (cl) Boltzmann statistics.

dtype: string

options: ['qn', 'cl']

default: qn

FIELDS

prefix

Prefix of the output files.

dtype: string

default:

asr

The method used to project out zero modes coming from continuous symmetries: crystal removes the three translational modes; molecule removes the three rotational modes in addition to the translational ones. none keeps all the modes.

dtype: string

options: ['none', 'crystal', 'poly']

default: none

random_type

Chooses the type of random numbers.

dtype: string

options: ['sobol', 'pseudo', 'file']

default: pseudo

displace_mode

The type of optimisation strategy for obtaining the mean position. sd stands for a steepest descent algorithm. ik stands for a Newton-Raphson scheme that requires the inverse of the force constant matrix \mathbf{K} . nmik stands for a Newton-Raphson scheme that only displaces along normal modes directions with statistically significant forces. rnmik same as nmik but performs several optimization steps using a reweighted sampling.

dtype: string

options: ['ik', 'sd', 'nmik', 'rnmik']

default: nmik

dynmat

The dynamical matrix of the trial Hamiltonian.

dtype: float

default: []

max_steps

Maximum number of Monte carlo steps per SCP iteration.

dtype: integer

max_iter

Maximum number of SCP iterations.

dtype: integer

default: 1

tau

Step size along the gradient for the sd displace_mode

dtype: float

default: 1.0

wthreshold

Threshold on minimum Boltzmann weights before more statistics must be accumulated.

dtype: float

default: 0.9

precheck

Flag for checking statistical significance of forces before optimisation of mean position.

dtype: boolean

default: True

checkweights

Flag for checking Boltzmann weights for whether more statistics are required.

dtype: boolean

default: True

chop

Threshold below which frequencies are set to zero.

dtype: float

default: 1e-09

nparallel

The number of Monte Carlo forces to be evaluated (in parallel) per i-PI step.

dtype: integer

default: 1

batch_weight_exponent

The exponent used to suppress low batch weights.

dtype: integer

default: 1

5.24 normalmodes

Fill in.

ATTRIBUTES

mode

The algorithm to be used: independent mode framework (imf) and vibrational self consistent field (vscf).

dtype: string

options: ['imf', 'vscf']

default: imf

FIELDS

prefix

Prefix of the output files.

dtype: string

default:

asr

Removes the zero frequency vibrational modes depending on the symmetry of the system for general polyatomic molecules, and periodic crystal structures.

dtype: string

options: ['none', 'poly', 'crystal']

default: none

dynamat

Portion of the dynamical matrix known to the current point in the calculation.

dtype: float

default: []

nprim

Number of primitive unit cells in the simulation cell.

dtype: float

default: 1.0

fnmrms

Fraction of harmonic RMS displacement used to sample along normal mode.

dtype: float

default: 1.0

nevib

Multiple of harm vibr energy up to which BO surface is sampled.

dtype: float

default: 25.0

nint

Integration points for Hamiltonian matrix elements.

dtype: integer

default: 101

pair_range

The range of pair combinations of normal modes to be considered.

dtype: integer

default: []

nbasis

Number of SHO states used as basis for anharmonic wvfn.

dtype: integer

default: 10

athresh

Convergence threshold for absolute error in vibr free energy per degree of freedom.

dtype: float

dimension: energy

default: 3.6749322e-06

ethresh

Convergence thresh for fractional error in vibr free energy.

dtype: float

default: 0.01

alpha

The fraction of mean field potential to mix with the result of the previous SCF iteration.

dtype: float

default: 1.0

nkbt

Threshold for $(e - e_{gs})/(k_B T)$ of vibr state to be incl in the VSCF and partition function.

dtype: float

default: 4.0

nexc

Minimum number of excited n-body states to calculate (also in MP2 correction).

dtype: integer

default: 5

mptwo

Flag determining whether MP2 correction is calculated.

dtype: boolean

default: False

solve

Flag determining whether the VSCF mean field Schroedinger's equation is solved.

dtype: boolean

default: False

grid

Flag determining whether the coupling potential is gridded or not.

dtype: boolean

default: True

print_mftpot

Flag determining whether MFT potentials are printed to file.

dtype: boolean

default: False

print_1b_map

Flag determining whether the independent mode potentials are printed to file.

dtype: boolean

default: False

print_2b_map

Flag determining whether the two body mapped coupling potentials are printed to file.

dtype: boolean

default: False

print_vib_density

Flag determining whether the vibrational density (ψ^2) are printed to file.

dtype: boolean

default: False

threebody

Flag determining whether three-mode coupling terms are accounted for.

dtype: boolean

default: False

nparallel

The number of forces evaluations per i-PI step.

dtype: integer

default: 1

5.25 optimizer

A Geometry Optimization class implementing most of the standard methods

ATTRIBUTES

mode

The geometry optimization algorithm to be used

dtype: string

options: ['sd', 'cg', 'bfgs', 'bfgstrm', 'lbfgs', 'damped_bfgs']

default: lbfgs

FIELDS

ls_options

“Options for line search methods. Includes: tolerance: stopping tolerance for the search (as a fraction of the overall energy tolerance), iter: the maximum number of iterations, step: initial step for bracketing, adaptive: whether to update initial step.

exit_on_convergence

Terminates the simulation when the convergence criteria are met.

dtype: boolean

default: True

tolerances

Convergence criteria for optimization. Default values are extremely conservative. Set them to appropriate values for production runs.

biggest_step

The maximum step size for (L)-BFGS line minimizations.

dtype: float

default: 100.0

scale_lbfgs

Scale choice for the initial hessian. 0 identity. 1 Use first member of position/gradient list. 2 Use last member of position/gradient list.

dtype: integer

default: 2

corrections_lbfgs

The number of past vectors to store for L-BFGS.

dtype: integer

default: 6

old_pos

The previous positions in an optimization step.

dtype: float

dimension: length

default: []

old_pot

The previous potential energy in an optimization step.

dtype: float

dimension: energy

default: []

old_force

The previous force in an optimization step.

dtype: float

dimension: force

default: []

old_direction

The previous direction in a CG or SD optimization.

dtype: float

default: []

invhessian_bfgs

Approximate inverse Hessian for BFGS, if known.

dtype: float

default: []

hessian_trm

Approximate Hessian for trm, if known.

dtype: float

default: []

tr_trm

The trust radius in trm.

dtype: float

dimension: length

default: []

qlist_lbfgs

List of previous position differences for L-BFGS, if known.

dtype: float

default: []

glist_lbfgs

List of previous gradient differences for L-BFGS, if known.

dtype: float

default: []

5.26 neb_optimizer

Contains the required parameters for performing nudged elastic band (NEB) calculations

ATTRIBUTES

mode

The geometry optimization algorithm to optimize NEB path

dtype: string

options: ['bfgstrm', 'damped_bfgs', 'fire']

default: fire

FIELDS

tolerances

Tolerance criteria to stop NEB optimization. If you work with DFT, do not use these defaults.

old_coord

The previous position in an optimization step.

dtype: float

dimension: length

default: []

full_force

The previous full-dimensional force in an optimization step.

dtype: float

dimension: force

default: []

full_pots

Previous physical potentials of all beads.

dtype: float

dimension: energy

default: []

old_nebpotential

Previous NEB potential energy, which includes spring energy.

dtype: float

default: []

old_nebgradient

The previous gradient including NEB spring forces.

dtype: float

default: []

old_direction

The previous direction.

dtype: float

default: []

biggest_step

The maximum atomic displacement in a single step of optimizations within NEB procedure. If requested step is larger, it will be downscaled so that maximal atomic displacement won't exceed biggest_step.

dtype: float

dimension: length

default: 0.5

scale_lbfgs

Scale choice for the initial hessian. 0 identity. 1 Use first member of position/gradient list. 2 Use last member of position/gradient list.

dtype: integer

default: 2

hessian_bfgs

Approximate Hessian for damped_BFGS, if known.

dtype: float

default: []

qlist_lbfgs

List of previous position differences for L-BFGS, if known.

dtype: float

default: []

glist_lbfgs

List of previous gradient differences for L-BFGS, if known.

dtype: float

default: []

corrections_lbfgs

The number of past vectors to store for L-BFGS.

dtype: integer

default: 5

dtmax_fire

Maximum time interval per step for FIRE.

dtype: float

default: 1.0

v_fire

Current velocity for FIRE

dtype: float

default: []

alpha_fire

velocity mixing factor for FIRE

dtype: float

default: 0.1

N_down_fire

consecutive steps in downhill direction for FIRE

dtype: integer

default: 0

N_up_fire

consecutive steps in uphill direction

dtype: integer

default: 0

dt_fire

time per step

dtype: float

default: 0.1

endpoints

Geometry optimization of endpoints (not implemented yet)

spring

Uniform or variable spring constants along the elastic band

tangent

How to calculate tangents: simple averaging from the original 1998 paper, or the improved tangent estimate from J. Chem. Phys. 113, 9978 (2000)

dtype: string

options: ['plain', 'improved']

default: improved

stage

Stage of the NEB pipeline: optimization of endpoints, NEB itself, climbing image

dtype: string

options: ['endpoints', 'neb', 'climb']

default: neb

use_climb

Use climbing image NEB or not

dtype: boolean

default: False

climb_bead

The index of the climbing bead.

dtype: integer

default: -1

5.27 string_optimizer

Contains the required parameters for performing string minimal energy path optimization.

ATTRIBUTES

mode

The geometry optimization algorithm to optimize MEP string

dtype: string

options: ['sd', 'cg', 'bfgs', 'bfgstrm', 'damped_bfgs', 'lbfgs', 'fire', 'euler']

default: bfgstrm

FIELDS

tolerances

Tolerance criteria to stop String optimization. If you work with DFT, do not use these defaults.

old_coord

The previous position in an optimization step.

dtype: float

dimension: length

default: []

full_force

The previous full-dimensional force in an optimization step.

dtype: float

dimension: force

default: []

full_pots

Previous physical potentials of all beads.

dtype: float

dimension: energy

default: []

old_stringpotential

Previous string potential energy.

dtype: float

default: []

old_stringgradient

The previous gradient of the string.

dtype: float

default: []

old_direction

The previous direction.

dtype: float

default: []

biggest_step

The maximum atomic displacement in a single step of optimizations within String MEP procedure. If requested step is larger, it will be downscaled so that maximal atomic displacement won't exceed biggest_step.

dtype: float

dimension: length

default: 0.5

scale_lbfgs

Scale choice for the initial hessian. 0 identity. 1 Use first member of position/gradient list. 2 Use last member of position/gradient list.

dtype: integer

default: 2

hessian_bfgs

Approximate Hessian for damped_BFGS, if known.

dtype: float

default: []

qlist_lbfgs

List of previous position differences for L-BFGS, if known.

dtype: float

default: []

glist_lbfgs

List of previous gradient differences for L-BFGS, if known.

dtype: float

default: []

corrections_lbfgs

The number of past vectors to store for L-BFGS.

dtype: integer

default: 5

tr_trm

Starting value for the trust radius for BFGSTRM.

dtype: float

dimension: length

default:

[1.]

dtmax_fire

Maximum time interval per step for FIRE.

dtype: float

default: 1.0

v_fire

Current velocity for FIRE

dtype: float

default: []

alpha_fire

velocity mixing factor for FIRE

dtype: float

default: 0.1

N_down_fire

consecutive steps in downhill direction for FIRE

dtype: integer

default: 0

N_up_fire

consecutive steps in uphill direction

dtype: integer

default: 0

dt_fire

time per step

dtype: float

default: 0.1

endpoints

Geometry optimization of endpoints (not implemented yet)

stage

Stage of the String pipeline: optimization of the endpoints, string opt., climbing image opt.

dtype: string

options: ['endpoints', 'string', 'climb']

default: string

use_climb

Use climbing image String MEP or not

dtype: boolean

default: False

climb_bead

The index of the climbing bead.

dtype: integer

default: -1

5.28 thermostat

Simulates an external heat bath to keep the velocity distribution at the correct temperature.

ATTRIBUTES

mode

The style of thermostating. 'langevin' specifies a white noise langevin equation to be attached to the cartesian representation of the momenta. 'svr' attaches a velocity rescaling thermostat to the cartesian representation of the momenta. Both 'pile_l' and 'pile_g' attaches a white noise langevin thermostat to the normal mode representation, with 'pile_l' attaching a local langevin thermostat to the centroid mode and 'pile_g' instead attaching a global velocity rescaling thermostat. 'gle' attaches a coloured noise langevin thermostat to the cartesian representation of the momenta, 'nm_gle' attaches a coloured noise langevin thermostat to the normal mode representation of the momenta and a langevin thermostat to the centroid and 'nm_gle_g' attaches a gle thermostat to the normal modes and a svr thermostat to the centroid. 'cl' represents a modified langevin thermostat which compensates for additional white noise from noisy forces or for dissipative effects. 'ffl' is the fast-forward langevin thermostat, in which momenta are flipped back whenever the action of the thermostat changes its direction. 'multiple' is a special thermostat mode, in which one can define multiple thermostats _inside_ the thermostat tag.

dtype: string

options: [' ', 'langevin', 'svr', 'pile_l', 'pile_g', 'gle', 'nm_gle', 'nm_gle_g', 'cl', 'ffl', 'multi']

FIELDS

ethermo

The initial value of the thermostat energy. Used when the simulation is restarted to guarantee continuity of the conserved quantity.

dtype: float

dimension: energy

default: 0.0

tau

The friction coefficient for white noise thermostats.

dtype: float

dimension: time

default: 0.0

pile_lambda

Scaling for the PILE damping relative to the critical damping. ($\gamma_k = 2 * \lambda * \omega_k$)

dtype: float

default: 1.0

pile_centroid_t

Option to set a different centroid temperature wrt. that of the ensemble. Only used if value other than 0.0.

dtype: float

dimension: temperature

default: 0.0

A

The friction matrix for GLE thermostats.

dtype: float

dimension: frequency

default: []

C

The covariance matrix for GLE thermostats.

dtype: float

dimension: temperature

default: []

s

Input values for the additional momenta in GLE.

dtype: float

dimension: ms-momentum

default: []

intau

The inherent noise time scale for compensating langevin thermostats.

dtype: float

dimension: time

default: 0.0

idtau

The inherent dissipation time scale for compensating langevin thermostats.

dtype: float

dimension: time

default: 0.0

apat

The time scale for automatic adjustment of CL thermostat's parameters.

dtype: float

dimension: time

default: 0.0

flip

Flipping type for ffl thermostat ('soft', 'hard', 'rescale', 'none')

dtype: string

default: rescale

thermostat

The thermostat for the atoms, keeps the atom velocity distribution at the correct temperature.

5.29 barostat

Simulates an external pressure bath.

ATTRIBUTES

mode

The type of barostat. 'isotropic' implements the Bussi-Zykova-Parrinello barostat [doi:10.1063/1.3073889] that isotropically scales the volume while sampling the isothermal isobaric ensemble. The implementation details are given in [doi:10.1016/j.cpc.2013.10.027]. 'sc-isotropic' implements the same for Suzuki-Chin path integral molecular dynamics [10.1021/acs.jctc.8b01297]. This barostat is suitable for simulating liquids. 'flexible' implements the path integral version of the Martyna-Tuckerman-Tobias-Klein barostat which incorporates full cell fluctuations while sampling the isothermal isobaric ensemble [doi:10.1063/1.478193]. This is suitable for anisotropic systems such as molecular solids. 'anisotropic' implements the Raiteri-Gale-Bussi barostat which enables cell fluctuations at constant external stress [10.1088/0953-8984/23/33/334213]. It is suitable for simulating solids at given external (non-diagonal) stresses and requires specifying a reference cell for estimating strain. Note that this ensemble is valid only within the elastic limit of small strains. For diagonal stresses (or external pressures) the 'flexible' and the 'anisotropic' modes should give very similar results. 'dummy' barostat does not do anything.

dtype: string

options: ['dummy', 'isotropic', 'flexible', 'anisotropic', 'sc-isotropic']

default: dummy

FIELDS

thermostat

The thermostat for the cell. Keeps the cell velocity distribution at the correct temperature. Note that the 'pile_l', 'pile_g', 'nm_gle' and 'nm_gle_g' options will not work for this thermostat.

tau

The time constant associated with the dynamics of the piston.

dtype: float

dimension: time

default: 1.0

p

Momentum (or momenta) of the piston.

dtype: float

dimension: momentum

default: []

h0

Reference cell for Parrinello-Rahman-like barostats. Should be roughly equal to the mean size of the cell averaged over the trajectory. Sampling might be inaccurate if the difference is too large.

dtype: float

dimension: length

default:

[0. 0. 0. 0. 0. 0. 0. 0.]

hfix

A list of the cell entries that should be held fixed (xx, yy, zz, xy, xz, yz). 'offdiagonal' is an alias for xy, xz, yz.

dtype: string

default: []

5.30 h0

Deals with the cell parameters. Takes as array which can be used to initialize the cell vector matrix.

dtype: float

dimension: length

ATTRIBUTES

units

The units the input data is given in.

dtype: string

default: automatic

shape

The shape of the array.

dtype: tuple

default: (0,)

mode

If 'mode' is 'manual', then the array is read in directly, then reshaped according to the 'shape' specified in a row-major manner. If 'mode' is 'file' then the array is read in from the file given.

dtype: string

options: ['manual', 'file']

default: manual

5.31 forcefield

Base forcefield class that deals with the assigning of force calculation jobs and collecting the data.

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.32 fsocket

Deals with the assigning of force calculation jobs to different driver codes, and collecting the data, using a socket for the data communication.

ATTRIBUTES

mode

Specifies whether the driver interface will listen onto a internet socket [inet] or onto a unix socket [unix].

dtype: string

options: ['unix', 'inet']

default: inet

matching

Specifies whether requests should be dispatched to any client, automatically matched to the same client when possible [auto] or strictly forced to match with the same client [lock].

dtype: string

options: ['auto', 'any', 'lock']

default: auto

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial. Should be set to True for FFSockets

dtype: boolean

default: True

FIELDS

address

This gives the server address that the socket will run on.

dtype: string

default: localhost

port

This gives the port number that defines the socket.

dtype: integer

default: 65535

slots

This gives the number of client codes that can queue at any one time.

dtype: integer

default: 4

exit_on_disconnect

Determines if i-PI should quit when a client disconnects.

dtype: boolean

default: False

timeout

This gives the number of seconds before assuming a calculation has died. If 0 there is no timeout.

dtype: float

default: 0.0

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

***default*:**

[-1]

5.33 fflj

Simple, internal LJ evaluator without cutoff, neighbour lists or minimal image convention. Expects standard LJ parameters, e.g. { eps: 0.1, sigma: 1.0 }.

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.34 ffdebye

Harmonic energy calculator

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

hessian

Specifies the Hessian of the harmonic potential. Default units are atomic. Units can be specified only by xml attribute. Implemented options are: 'atomic_unit', 'ev/ang^2'

dtype: float

dimension: hessian

default: []

x_reference

Minimum-energy configuration for the harmonic potential

dtype: float

dimension: length

default: []

v_reference

Zero-value of energy for the harmonic potential

dtype: float

dimension: energy

default: 0.0

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

***default*:**

[-1]

5.35 ffplumed

Direct PLUMED interface

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

file

This describes the location to read the reference structure file from.

dtype: string

default:

plumeddat

The PLUMED input file

dtype: string

default: plumed.dat

plumedstep

The current step counter for PLUMED calls

dtype: integer

default: 0

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.36 ffyaff

Uses a Yaff force field to compute the forces.

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

yaffpara

This gives the file name of the Yaff input parameter file.

dtype: string

default: parameters.txt

yaffsys

This gives the file name of the Yaff input system file.

dtype: string

default: system.chk

yafflog

This gives the file name of the Yaff output log file.

dtype: string

default: yaff.log

rcut

This gives the real space cutoff used by all pair potentials in atomic units.

dtype: float

default: 18.89726133921252

alpha_scale

This gives the alpha parameter in the Ewald summation based on the real-space cutoff: $\alpha = \alpha_scale / rcut$. Higher values for this parameter imply a faster convergence of the reciprocal terms, but a slower convergence in real-space.

dtype: float

default: 3.5

gcut_scale

This gives the reciprocal space cutoff based on the alpha parameter: $gcut = gcut_scale * \alpha$. Higher values for this parameter imply a better convergence in the reciprocal space.

dtype: float

default: 1.1

skin

This gives the skin parameter for the neighborlist.

dtype: integer

default: 0

smooth_ei

This gives the flag for smooth truncations for the electrostatic interactions.

dtype: boolean

default: False

reci_ei

This gives the method to be used for the reciprocal contribution to the electrostatic interactions in the case of periodic systems. This must be one of 'ignore' or 'ewald'. The 'ewald' option is only supported for 3D periodic systems.

dtype: string

default: ewald

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

***default*:**

[-1]

5.37 ffsgdml

A SGDML energy calculator

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

sGDML_model

This gives the file name of the sGDML model.

dtype: string

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.38 ffdmd

Simple, internal DMD evaluator without neighbor lists, but with PBC. Expects coupling elements ($n*(n-1)/2$ of them), oscillating frequency and time step.

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

dmd_coupling

Specifies the coupling between atom pairs (should be size $N*(N-1)/2$ ordered c21, c32, c31, c43, c42, c41 etc. – in atomic units!)

dtype: float

default: []

dmd_freq

Frequency of the oscillation of the time-dependent term

dtype: float

dimension: frequency

default: 0.0

dmd_dt

Time step of the oscillating potential. Should match time step of simulation

dtype: float

dimension: time

default: 0.0

dmd_step

The current step counter for dmd.

dtype: integer

default: 0

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.39 fcommittee

Combines multiple forcefields to build a committee model, that can be used to compute uncertainty-quantified machine-learning models. Each forcefield can be any of the other FF objects, and each should be used with a client that generates a slightly different estimation of energy and forces. These are averaged, and the mean used as the actual forcefield. Statistics about the distribution are also returned as extras fields, and can be printed for further postprocessing. Also contains options to use it for uncertainty estimation and for active learning in a ML context, based on a committee model. Implements the approaches discussed in DOI: 10.1063/5.0036522.

ATTRIBUTES

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial

dtype: boolean

default: False

FIELDS

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

weights

List of weights to be given to the forcefields. Defaults to 1 for each FF. Note that the components are divided by the number of FF, and so the default corresponds to an average.

dtype: float

default: []

alpha

Scaling of the variance of the model, corresponding to a calibration of the error

dtype: float

default: 1.0

baseline_name

Name of the forcefield object that should be treated as the baseline for a weighted baseline model.

dtype: string

default:

baseline_uncertainty

Corresponds to the expected error of the baseline model. This represents the error on the TOTAL potential energy of the simulation.

dtype: float

dimension: energy

default: -1.0

active_thresh

The uncertainty threshold for active learning. Structure with an uncertainty above this value are printed in the specified output file so they can be used for active learning.

dtype: float

dimension: energy

default: 0.0

active_output

Output filename for structures that exceed the accuracy threshold of the model, to be used in active learning.

dtype: string

default: active_output

ffsocket

Deals with the assigning of force calculation jobs to different driver codes, and collecting the data, using a socket for the data communication.

fflj

Simple, internal LJ evaluator without cutoff, neighbour lists or minimal image convention. Expects standard LJ parameters, e.g. { eps: 0.1, sigma: 1.0 }.

ffdebye

Harmonic energy calculator

ffplumed

Direct PLUMED interface

ffyaaff

Uses a Yaff force field to compute the forces.

ffsgdml

A SGDML energy calculator

5.40 ff cavphsocket

A cavity molecular dynamics driver for vibrational strong coupling. In the current implementation, only a single cavity mode polarized along the x and y directions is coupled to the molecules. Check <https://doi.org/10.1073/pnas.2009272117> and also `examples/lammps/h2o-cavmd/` for details.

ATTRIBUTES

mode

Specifies whether the driver interface will listen onto a internet socket [inet] or onto a unix socket [unix].

dtype: string

options: ['unix', 'inet']

default: inet

matching

Specifies whether requests should be dispatched to any client, automatically matched to the same client when possible [auto] or strictly forced to match with the same client [lock].

dtype: string

options: ['auto', 'any', 'lock']

default: auto

name

Mandatory. The name by which the forcefield will be identified in the System forces section.

dtype: string

pbc

Applies periodic boundary conditions to the atoms coordinates before passing them on to the driver code.

dtype: boolean

default: True

threaded

Whether the forcefield should use a thread loop to evaluate, or work in serial. Should be set to True for FFSockets

dtype: boolean

default: True

FIELDS

charge_array

The partial charges of all the atoms, in the format [Q1, Q2, ...].

dtype: float

dimension: length

default: []

apply_photon

Determines if additional photonic degrees of freedom is included or not.

dtype: boolean

default: False

E0

The value of varepsilon (effective light-matter coupling strength) in atomic units.

dtype: float

default: 0.0

omega_c

This gives the cavity photon frequency at normal incidence.

dtype: float

dimension: frequency

default: 0.01

ph_rep

In the current implementation, two energy-degenerate photon modes polarized along x and y directions are coupled to the molecular system. If 'loose', the cavity photons polarized along the x, y directions are represented by two 'L' atoms; the x dimension of the first 'L' atom is coupled to the molecules, and the y dimension of the second 'L' atom is coupled to the molecules. If 'dense', the cavity photons polarized along the x, y directions are represented by one 'L' atom; the x and y dimensions of this 'L' atom are coupled to the molecules.

dtype: string

options: ['loose', 'dense']

default: loose

address

This gives the server address that the socket will run on.

dtype: string

default: localhost

port

This gives the port number that defines the socket.

dtype: integer

default: 65535

slots

This gives the number of client codes that can queue at any one time.

dtype: integer

default: 4

exit_on_disconnect

Determines if i-PI should quit when a client disconnects.

dtype: boolean

default: False

timeout

This gives the number of seconds before assuming a calculation has died. If 0 there is no timeout.

dtype: float

default: 0.0

latency

The number of seconds the polling thread will wait between examining the list of requests.

dtype: float

default: 0.0001

parameters

The parameters of the force field

dtype: dictionary

default: { }

activelist

List with indexes of the atoms that this socket is taking care of. Default: [-1] (corresponding to all)

dtype: integer

default:

[-1]

5.41 forces

Deals with creating all the necessary forcefield objects.

FIELDS

force

The class that deals with how each forcefield contributes to the overall potential, force and virial calculation.

5.42 force

The class that deals with how each forcefield contributes to the overall potential, force and virial calculation.

ATTRIBUTES

nbeads

If the forcefield is to be evaluated on a contracted ring polymer, this gives the number of beads that are used. If not specified, the forcefield will be evaluated on the full ring polymer.

dtype: integer

default: 0

weight

A scaling factor for this forcefield, to be applied before adding the force calculated by this forcefield to the total force.

dtype: float

default: 1.0

fd_epsilon

The finite displacement to be used for calculating the Suzuki-Chin contribution of the force. If the value is negative, a centered finite-difference scheme will be used. [in bohr]

dtype: float

default: -0.001

name

An optional name to refer to this force component.

dtype: string

default:

forcefield

Mandatory. The name of the forcefield this force is referring to.

dtype: string

default:

FIELDS

mts_weights

The weight of force in each mts level starting from outer.

dtype: float

dimension: force

default:

[1.]

force_extras

A list of quantities that should be extracted from the 'extra' string returned by the forcefield, that are then treated the same way as energy or forces – that is treated as a numerical, physical quantity, interpolated when changing the number of PI replicas. Same quantities from different force components are summed as well. The names should correspond to entries in the JSON-formatted extra string.

dtype: string

default: []

5.43 al6xxx_kmc

Holds all the information for the KMC dynamics, such as timestep, rates and barriers that control it.

ATTRIBUTES

mode

The KMC algorithm to be used

dtype: string

options: ['rfkmc']

default: rfkmc

FIELDS

optimizer

Option for geometry optimization step

nstep

The number of optimization steps.

dtype: integer

default: 10

a0

FCC lattice parameter

dtype: float

dimension: length

default: 1.0

diffusion_barrier_al

Barrier for vacancy diffusion in pure Al.

dtype: float

dimension: energy

default: 0.01

diffusion_prefactor_al

Prefactor for vacancy diffusion in pure Al.

dtype: float

dimension: frequency

default: 2.4188843e-05

diffusion_barrier_mg

Barrier for vacancy-assisted diffusion of Mg.

dtype: float

dimension: energy

default: 0.0

diffusion_prefactor_mg

Prefactor for vacancy-assisted diffusion of Mg.

dtype: float

dimension: frequency

default: 0.0

diffusion_barrier_si

Barrier for vacancy-assisted diffusion of Si.

dtype: float

dimension: energy

default: 0.0

diffusion_prefactor_si

Prefactor for vacancy-assisted diffusion of Si.

dtype: float

dimension: frequency

default: 0.0

neval

The number of parallel force evaluators.

dtype: integer

default: 4

ncell

The number of repeat cells in each direction.

dtype: integer

default: 4

nvac

The number of vacancies.

dtype: integer

default: 4

nsi

The number of silicon atoms.

dtype: integer

default: 4

nmg

The number of magnesium atoms.

dtype: integer

default: 4

idx

The position of the atoms on the lattice, relative to the canonical ordering.

dtype: integer

default: []

tottime

Total KMC time elapsed

dtype: float

dimension: time

default: 0.0

ecache_file

Filename for storing/loading energy cache

dtype: string

default:

qcache_file

Filename for storing/loading positions cache

dtype: string

default:

max_cache_len

Maximum cache length before oldest entry is deleted

dtype: integer

default: 1000

5.44 atoms

Deals with a single replica of the system or classical simulations.

FIELDS

natoms

The number of atoms.

dtype: integer

default: 0

q

The positions of the atoms, in the format [x1, y1, z1, x2, ...].

dtype: float

dimension: length

default: []

p

The momenta of the atoms, in the format [px1, py1, pz1, px2, ...].

dtype: float

dimension: momentum

default: []

m

The masses of the atoms, in the format [m1, m2, ...].

dtype: float

dimension: mass

default: []

names

The names of the atoms, in the format [name1, name2, ...].

dtype: string

default: []

5.45 beads

Describes the bead configurations in a path integral simulation.

ATTRIBUTES

natoms

The number of atoms.

dtype: integer

default: 0

nbeads

The number of beads.

dtype: integer

default: 0

FIELDS

q

The positions of the beads. In an array of size [nbeads, 3*natoms].

dtype: float

dimension: length

default: []

p

The momenta of the beads. In an array of size [nbeads, 3*natoms].

dtype: float

dimension: momentum

default: []

m

The masses of the atoms, in the format [m1, m2, ...].

dtype: float

dimension: mass

default: []

names

The names of the atoms, in the format [name1, name2, ...].

dtype: string

default: []

5.46 prng

Deals with the pseudo-random number generator.

FIELDS

seed

This is the seed number used to generate the initial state of the random number generator.

dtype: integer

default: 123456

state

Gives the state vector for the random number generator. Avoid directly modifying this unless you are very familiar with the inner workings of the algorithm used.

dtype: integer

default: []

has_gauss

Determines whether there is a stored gaussian number or not. A value of 0 means there is none stored.

dtype: integer

default: 0

gauss

The stored Gaussian number.

dtype: float

default: 0.0

set_pos

Gives the position in the state array that the random number generator is reading from.

dtype: integer

default: 0

5.47 normal_modes

Deals with the normal mode transformations, including the adjustment of bead masses to give the desired ring polymer normal mode frequencies if appropriate. Takes as arguments frequencies, of which different numbers must be specified and which are used to scale the normal mode frequencies in different ways depending on which 'mode' is specified.

ATTRIBUTES

transform

Specifies whether to calculate the normal mode transform using a fast Fourier transform or a matrix multiplication. For small numbers of beads the matrix multiplication may be faster.

dtype: string

options: ['fft', 'matrix']

default: fft

propagator

How to propagate the free ring polymer dynamics. Cayley transform is not exact but is strongly stable and avoid potential resonance issues. A bab scheme performs numerical verlet type propagation. All three options work for distinguishable particles. Only the bab propagator can be used with bosonic particles.

dtype: string

options: ['exact', 'cayley', 'bab']

default: exact

FIELDS

frequencies

Specifies normal mode frequencies for a (closed path) calculation

dtype: float

dimension: frequency

default: []

open_paths

Indices of the atoms whose path should be opened (zero-based).

dtype: integer

default: []

bosons

Indices of the atoms that are bosons (zero-based).

dtype: integer

default: []

nmts

The number of iterations to perform one bab step.

dtype: integer

default: 10

5.48 frequencies

Provides a compact way of specifying the ring polymer frequencies

dtype: float

dimension: frequency

ATTRIBUTES

units

The units the input data is given in.

dtype: string

default: automatic

shape

The shape of the array.

dtype: tuple

default: (0,)

mode

If 'mode' is 'manual', then the array is read in directly, then reshaped according to the 'shape' specified in a row-major manner. If 'mode' is 'file' then the array is read in from the file given.

dtype: string

options: ['manual', 'file']

default: manual

style

Specifies the technique to be used to calculate the dynamical masses. 'rpmd' simply assigns the bead masses the physical mass. 'manual' sets all the normal mode frequencies except the centroid normal mode manually. 'pa-cmd' takes an argument giving the frequency to set all the non-centroid normal modes to. 'wmax-cmd' is similar to 'pa-cmd', except instead of taking one argument it takes two ([wmax,wtarget]). The lowest-lying normal mode will be set to wtarget for a free particle, and all the normal modes will coincide at frequency wmax.

dtype: string

options: ['pa-cmd', 'wmax-cmd', 'manual', 'rpmd']

default: rpmd

5.49 initialize

Specifies the number of beads, and how the system should be initialized.

ATTRIBUTES

nbeads

The number of beads. Will override any provision from inside the initializer. A ring polymer contraction scheme is used to scale down the number of beads if required. If instead the number of beads is scaled up, higher normal modes will be initialized to zero.

dtype: integer

FIELDS

positions

Initializes atomic positions. Will take a 'units' attribute of dimension 'length'

dtype: string

velocities

Initializes atomic velocities. Will take a 'units' attribute of dimension 'velocity'

dtype: string

momenta

Initializes atomic momenta. Will take a 'units' attribute of dimension 'momentum'

dtype: string

masses

Initializes atomic masses. Will take a 'units' attribute of dimension 'mass'

dtype: string

labels

Initializes atomic labels

dtype: string

cell

Initializes the configuration of the cell. Will take a 'units' attribute of dimension 'length'

dtype: string

file

Initializes everything possible for the given mode. Will take a 'units' attribute of dimension 'length'. The unit conversion will only be applied to the positions and cell parameters. The 'units' attribute is deprecated. Append a 'quantity{units}' to the comment line of the xyz or to the 'TITLE' tag of a pdb.

dtype: string

gle

Initializes the additional momenta in a GLE thermostat.

dtype: string

5.50 file

This is the class to initialize from file.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'chk' stands for initialization from a checkpoint file. 'ase' is to read a file with the Atomic Simulation Environment

dtype: string

options: ['xyz', 'pdb', 'chk', 'ase']

default: chk

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

cell_units

The units for the cell dimensions.

dtype: string

default: automatic

5.51 positions

This is the class to initialize positions.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'ase' is to read a file with the Atomic Simulation Environment. 'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector.

dtype: string

options: ['manual', 'xyz', 'pdb', 'ase', 'chk']

default: chk

index

The index of the atom for which the value will be set. If a negative value is specified, then all atoms are assumed.

dtype: integer

default: -1

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

5.52 momenta

This is the class to initialize momenta.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector. 'thermal' means that the data is to be generated from a Maxwell-Boltzmann distribution at the given temperature.

dtype: string

options: ['manual', 'xyz', 'pdb', 'ase', 'chk', 'thermal']

default: chk

index

The index of the atom for which the value will be set. If a negative value is specified, then all atoms are assumed.

dtype: integer

default: -1

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

5.53 labels

This is the class to initialize atomic labels.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'ase' is to read a file with the Atomic Simulation Environment. 'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector.

dtype: string

options: ['manual', 'xyz', 'pdb', 'ase', 'chk']

default: chk

index

The index of the atom for which the value will be set. If a negative value is specified, then all atoms are assumed.

dtype: integer

default: -1

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

5.54 masses

This is the class to initialize atomic masses.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'ase' is to read a file with the Atomic Simulation Environment. 'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector.

dtype: string

options: ['manual', 'xyz', 'pdb', 'ase', 'chk']

default: chk

index

The index of the atom for which the value will be set. If a negative value is specified, then all atoms are assumed.

dtype: integer

default: -1

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

5.55 velocities

This is the class to initialize velocities.

dtype: string

ATTRIBUTES

mode

The input data format. 'xyz' and 'pdb' stand for xyz and pdb input files respectively. 'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector. 'thermal' means that the data is to be generated from a Maxwell-Boltzmann distribution at the given temperature.

dtype: string

options: ['manual', 'xyz', 'pdb', 'ase', 'chk', 'thermal']

default: chk

index

The index of the atom for which the value will be set. If a negative value is specified, then all atoms are assumed.

dtype: integer

default: -1

bead

The index of the bead for which the value will be set. If a negative value is specified, then all beads are assumed.

dtype: integer

default: -1

5.56 init_cell

This is the class to initialize cell.

dtype: string

ATTRIBUTES

mode

This decides whether the system box is created from a cell parameter matrix, or from the side lengths and angles between them. If 'mode' is 'manual', then 'cell' takes a 9-elements vector containing the cell matrix (row-major). The 1st element define lattice vector a, the 2nd, 5th elements define lattice vector b, and the 3rd, 6th, 9th elements define lattice vector c. The other elements are ignored. If 'mode' is 'abcABC', then 'cell' takes an array of 6 floats, the first three being the length of the sides of the system parallelepiped, and the last three being the angles (in degrees) between those sides. Angle A corresponds to the angle between sides b and c, and so on for B and C. If mode is 'abc', then this is the same as for 'abcABC', but the cell is assumed to be orthorhombic. 'pdb' and 'chk' read the cell from a PDB or a checkpoint file, respectively.

dtype: string

options: ['manual', 'pdb', 'chk', 'abc', 'abcABC']

default: manual

5.57 gle

This is the class to initialize the thermostat (ethermo and fictitious momenta).

dtype: string

ATTRIBUTES

mode

'chk' stands for initialization from a checkpoint file. 'manual' means that the value to initialize from is giving explicitly as a vector.

dtype: string

options: ['chk', 'manual']

default: manual

5.58 output

This class defines how properties, trajectories and checkpoints should be output during the simulation. May contain zero, one or many instances of properties, trajectory or checkpoint tags, each giving instructions on how one output file should be created and managed.

ATTRIBUTES

prefix

A string that will be prepended to each output file name. The file name is given by 'prefix'.filename' + format_specifier. The format specifier may also include a number if multiple similar files are output.

dtype: string

default: i-pi

FIELDS

properties

Each of the properties tags specify how to create a file in which one or more properties are written, one line per frame.

dtype: string

trajectory

Each of the trajectory tags specify how to create a trajectory file, containing a list of per-atom coordinate properties.

dtype: string

checkpoint

Each of the checkpoint tags specify how to create a checkpoint file, which can be used to restart a simulation.

dtype: integer

5.59 properties

This class deals with the output of properties to one file. Between each property tag there should be an array of strings, each of which specifies one property to be output.

dtype: string

ATTRIBUTES

shape

The shape of the array.

dtype: tuple

default: (0,)

mode

If 'mode' is 'manual', then the array is read in directly, then reshaped according to the 'shape' specified in a row-major manner. If 'mode' is 'file' then the array is read in from the file given.

dtype: string

options: ['manual', 'file']

default: manual

filename

A string to specify the name of the file that is output. The file name is given by 'prefix'.filename' + format_specifier. The format specifier may also include a number if multiple similar files are output.

dtype: string

default: out

stride

The number of steps between successive writes.

dtype: integer

default: 1

flush

How often should streams be flushed. 1 means each time, zero means never.

dtype: integer

default: 1

5.60 checkpoint

This class defines how a checkpoint file should be output. Optionally, between the checkpoint tags, you can specify one integer giving the current step of the simulation. By default this integer will be zero.

dtype: integer

ATTRIBUTES

filename

A string to specify the name of the file that is output. The file name is given by 'prefix'.filename' + format_specifier. The format specifier may also include a number if multiple similar files are output.

dtype: string

default: restart

stride

The number of steps between successive writes.

dtype: integer

default: 1

overwrite

This specifies whether or not each consecutive checkpoint file will overwrite the old one.

dtype: boolean

default: True

5.61 trajectory

This class defines how one trajectory file should be output. Between each trajectory tag one string should be given, which specifies what data is to be output.

dtype: string

ATTRIBUTES

filename

A string to specify the name of the file that is output. The file name is given by 'prefix'.filename' + format_specifier. The format specifier may also include a number if multiple similar files are output.

dtype: string

default: traj

stride

The number of steps between successive writes.

dtype: integer

default: 1

format

The output file format.

dtype: string

options: ['xyz', 'pdb', 'ase']

default: xyz

cell_units

The units for the cell dimensions.

dtype: string

default:

bead

Print out only the specified bead. A negative value means print only one every -(bead) beads, e.g. -2 means print just the even beads, -4 one every four and so on.

dtype: integer

default: -1

flush

How often should streams be flushed. 1 means each time, zero means never.

dtype: integer

default: 1

extra_type

What extra to print from the different extra strings.

dtype: string

default: raw

6 Internal units

All the units used internally by i-PI are atomic units, as given below. By default, both input and output data are given in atomic units, but in most cases the default units can be overridden if one wishes so. For details on how to do this, see 3.1.1.1 and 3.2.1.

Unit	Name	S.I. Value
Length	Bohr radius	5.2917721e-11 m
Time	N.A.	2.4188843e-17 s
Mass	Electron mass	9.1093819e-31 kg
Temperature	Hartree	315774.66 K
Energy	Hartree	4.3597438e-18 J
Pressure	N.A.	2.9421912e13 Pa

Regarding the specification of these units in the i-PI input files, the user is able to specify units both in the i-PI input file or in the structure file. If the structure file is of the .xyz format, the units specifications should be present in the comment line. Examples of such inputs can be found in `examples/pes-regtest/io-units/`. The code then behaves in the following way, depending on the user's choice:

- If no units are specified in the input, i-PI tries to guess from the comment line of the structure file and if nothing is present, assumes atomic units, except for the cases discussed below.
- If units are specified in both input and structure file and they match, conversion happens just once. If they do not match, an error is raised and i-PI stops.

There are a few exceptions:

- The total simulation time, specified by the flag `<total_time>`, has to be provided in seconds.
- The PDB standard defines units to be in angstrom, and so default units for PDB are angstrom.
- When using ASE format for I/O, the units are dictated by the ASE defaults. So for example, the ASE extended xyz format expects positions in angstroms. (See for example `examples/ase-io/`)

7 Troubleshooting

7.1 Input errors

- *not well-formed (invalid token)*: Seen if the input file does not have the correct xml syntax. Should be accompanied by a line number giving the point in the file where the syntax is incorrect.
- *mismatched tag*: One of the closing tags does not have the same name as the corresponding opening tag. Could be caused either by a misspelling of one of the tags, or by having the closing tag in the wrong place. This last one is a standard part of the xml syntax, if the opening tag of one item is after the opening tag of a second, then its closing tag should be before the closing tag of the second. Should be accompanied by a line number giving the position of the closing tag.
- *Uninitialized value of type __ or Attribute/Field name __ is mandatory and was not found in the input for property __*: The xml file is missing a mandatory tag, i.e. one without which the simulation cannot be initialized. Find which tag name is missing and add it.

- *Attribute/tag name ___ is not a recognized property of ___ objects*: The first tag should not be found within the second set of tags. Check that the first tag is spelt correctly, and that it has been put in the right place.
- *___ is not a valid option (___)*: This attribute/tag only allows a certain range of inputs. Pick one of the items from the list given.
- *___ is an undefined unit for kind ___ or ___ is not a valid unit prefix or Unit ___ is not structured with a prefix+base syntax*: The unit input by the user is not correct. Make sure it corresponds to the correct dimensionality, and is spelt correctly.
- *Invalid literal for int() with base 10: ___ or Invalid literal for float(): ___ or ___ does not represent a bool value*: The data input by the user does not have the correct data type. See section 3.1.1 for what constitutes a valid integer/float/boolean value.
- *Error in list syntax: could not locate delimiters*: The array input data did not have the required braces. For a normal array use [], for a dictionary use { }, and for a tuple use ().
- *The number of atom records does not match the header of xyz file*: Self-explanatory.
- *list index out of range*: This will normally occur if the configuration is initialized from an invalid input file. This will either cause the code to try to read part of the input file that does not exist, or to set the number of beads to zero which causes this error in a different place. Check that the input file has the correct syntax.

7.2 Initialization errors

- *Negative ___ parameter specified.*: Self-explanatory.
- *If you are initializing cell from cell side lengths you must pass the 'cell' tag an array of 3 floats*: If you are attempting to initialize a cell using the “abc” mode, the code expects three floats corresponding to the three side lengths.
- *If you are initializing cell from cell side lengths and angles you must pass the 'cell' tag an array of 6 floats*: If you are attempting to initialize a cell using the “abcABC” mode, the code expects six floats corresponding to the three side lengths, followed by the three angles in degrees.
- *Cell objects must contain a 3x3 matrix describing the cell vectors.*: If you are attempting to initialize a cell using the “manual” mode, the code expects nine floats corresponding to the cell vector matrix side lengths. Note that the values of the lower-diagonal elements will be set to zero.
- *Array shape mismatch in q/p/m/names in beads input*: The size of the array in question does not have the correct number of elements given the number of atoms and the number of beads used in the rest of the input. If the number of beads is nbeads and the number of atoms natoms, then q and p should have shape (nbeads, 3*natoms) and m and names should have shape (natoms,).
- *No thermostat/barostat tag provided for NVT/NPT simulation*: Some ensembles can only be sampled if a thermostat and/or barostat have been defined, and so for simulations at constant temperature and/or pressure these tags are mandatory. If you wish to not use a thermostat/barostat, but still want to keep the ensemble the same, then use “dummy” mode thermostat/barostat, which simply does nothing.
- *Pressure/Temperature should be supplied for constant pressure/temperature simulation*: Since in this case the ensemble is defined by these parameters, these must be input by the user. Add the appropriate tags to the input file.
- *Manual path mode requires (nbeads-1) frequencies, one for each internal mode of the path.*: If the “mode” tag of “normal_modes” is set to “manual”, it will expect an array of frequencies, one for each of the internal normal modes of the ring polymers.
- *PA-CMD mode requires the target frequency of all the internal modes.*: If the “mode” tag of “normal_modes” is set to “pa-cmd”, it will expect an array of one frequency, to which all the internal modes of the ring polymers will be set.

- *WMAX-CMD mode requires [wmax, wtarget]. The normal modes will be scaled such that the first internal mode is at frequency wtarget and all the normal modes coincide at frequency wmax.*: If the “mode” tag of “normal_modes” is set to “wmax-cmd”, it will expect an array of two frequencies, one to set the lowest frequency normal mode, and one for the other normal mode frequencies.
- *Number of beads __ doesn’t match GLE parameter nb=__*: The matrices used to define the generalized Langevin equations of motion do not have the correct first dimension. If matrices have been downloaded from <http://imx-cosmo.github.io/gle4md/> make sure that you have input the correct number of beads.
- *Initialization tries to match up structures with different atom numbers*: If in the initialization any of the matrices has an array shape which do not correspond to the same number of atoms, then they cannot correspond to the same system. Check the size of the arrays specified if they have been input manually.
- *Cannot initialize single atom/bead as atom/bead index __ is larger than the number of atoms/beads*: Self-explanatory. However, note that indices are counted from 0, so the first replica/atom is defined by an index 0, the second by an index 1, and so on.
- *Cannot initialize the momenta/masses/labels/single atoms before the size of the system is known.*: In the code, a beads object is created to hold all the information related to the configuration of the system. However, until a position vector has been defined, this object is not created. Therefore, whichever arrays are being initialized individually, the position vector must always be initialized first.
- *Trying to resample velocities before having masses.*: A Maxwell-Boltzmann distribution is partly defined by the atomic masses, and so the masses must be defined before the velocities can be resampled from this distribution.
- *Cannot thermalize a single bead*: It does not make sense to initialize the momenta of only one of the beads, and so i-PI does not give this functionality.
- *Initializer could not initialize __*: A property of the system that is mandatory to properly run the simulation has not been initialized in either the “initialize” section or the appropriate section in beads.
- *Ensemble does not have a thermostat to initialize or There is nothing to initialize in non-GLE thermostats or Checkpoint file does not contain usable thermostat data*: These are raised if the user has tried to initialize the matrices for the GLE thermostats with a checkpoint file that either does not have a GLE thermostat or does not have a thermostat at all.
- *Size mismatch in thermostat initialization data*: Called if the shape of the GLE matrices defined in the checkpoint file is different from those defined in the new simulation.
- *Replay can only read from PDB or XYZ files – or a single frame from a CHK file*: If the user specifies a replay ensemble, the state of the system must be defined by either a configuration file or a checkpoint file, and cannot be specified manually.

7.3 Output errors

- *The stride length for the __ file output must be positive.*: Self-explanatory
- *__ is not a recognized property/output trajectory*: The string as defined in the “properties”/“trajectory” tag does not correspond to one of the available trajectories. Make sure that both the syntax is correct, and that the property has been spelt correctly.
- *Could not open file __ for output*: Raised if there is a problem opening the file defined by the “filename” attribute.
- *Selected bead index __ does not exist for trajectory __*: You have asked for the trajectory of a bead index greater than the number of the replicas of the system. Note that indices are counted from 0, so the first replica is defined by an index 0, the second by an index 1, and so on.
- *Incorrect format in unit specification __*: Usually raised if one of the curly braces has been neglected.

- *Incorrect format in argument list __*: This will be raised either if one of the brackets has been neglected, or if the delimiters between arguments, in this case “;”, are not correct. This is usually raised if, instead of separating the arguments using “;”, they are instead separated by “,”, since this causes the property array to be parsed incorrectly.
- *__ got an unexpected keyword argument __*: This will occur if one of the argument lists of one of the properties specified by the user has a keyword argument that does not match any of those in the function to calculate it. Check the properties.py module to see which property this function is calculating, and what the correct keyword arguments are. Then check the “properties” tag, and find which of the arguments has been misspelt.
- *Must specify the index of atom_vec property*: Any property which prints out a vector corresponding to one atom needs the index of that atom, as no default is specified.
- *Cannot output __ as atom/bead index __ is larger than the number of atoms/beads*: Self-explanatory. However, note that indices are counted from 0, so the first replica/atom is defined by an index 0, the second by an index 1, and so on.
- *Couldn't find an atom that matched the argument of __*: For certain properties, you can specify an atom index or label, so that the property is averaged only over the atoms that match it. If however no atom labels match the argument given, then the average will be undefined. Note that for properties which are cumulatively counted rather than averaged, this error is not raised, and if no atom matches the label given 0 will be returned.

7.4 Socket errors

- *Address already in use*: This is called if the server socket is already being used by the host network. There are several possible reasons for getting this error. Firstly, it might simply be that two simulations are running concurrently using the same host and port number. In this case simply change the port number of one of the simulations. Secondly, you can get this error if you try to rerun a simulation that previously threw an exception, since it takes a minute or so before the host will disconnect the server socket if it is not shut down cleanly. In this case, simply wait for it to disconnect, and try again. Finally, you will get this error if you try to use a restricted port number (i.e. below 1024) while not root. You should always use a non-restricted port number for i-PI simulations.
- *Error opening unix socket. Check if a file /tmp/ipi__ exists, and remove it if unused.*: Similar to the above error, but given if you are using a unix socket rather than an internet socket. Since this binds locally the socket can be removed by the user, which means that it is not necessary to wait for the computer to automatically disconnect an unused server socket.
- *Port number __ out of acceptable range*: The port number must be between 1 and 65535, and should be greater than 1024. Change the port number accordingly.
- *Slot number __ out of acceptable range*: The slot number must be between 1 and 5. Change the slot number accordingly.
- *'NoneType' object has no attribute 'Up'*: This is called if an exception is raised during writing the data to output, and so the thread that deals with the socket is terminated uncleanly. Check the stack trace for the original exception, since this will be the actual source of the problem. Also note that, since the socket thread was not cleaned up correctly, the server socket may not have been disconnected properly and you may have to wait for a minute before you can restart a simulation using the same host and port number.

7.5 Mathematical errors

- *math domain error*: If the cell parameters are defined using the side lengths and angles, with either a pdb file or using the “abcABC” initialization mode, then for some value of the angles it is impossible to construct a valid cell vector matrix. This will cause the code to attempt to take the square root of a negative number, which gives this exception.
- *overflow encountered in exp*: Sometimes occurs in *NPT* runs when the simulation box “explodes”. Make sure you have properly equilibrated the system before starting and that the timestep is short enough to not introduce very large integration errors.

7.6 I-PI is slow!

i-PI is not designed to be highly efficient, but it should not be the bottleneck in your calculations. Most of the time, if you see a major slow-down, there is a problem with your setup. Here are some tricks you can try.

- *use a UNIX domain socket*: if you run on a single node, it is much faster to use *mode=“unix”* in your `<ffsocket>` classes. This uses a shared-memory communication process, that avoids much of the network latency of a TCP/IP socket.
- *reduce latency*: if you have a VERY fast forcefield (with force evaluation below ~1ms) it might help to set the `<latency>` parameter of `<ffsocket>` to a small value, 1e-4s or less
- *reduce I/O*: outputting hundreds of beads configurations at each time step is going to be slow in any scenario, but particularly so when using text files and Python. Reduce the output frequency using a larger *stride*, and/or *flush* less often.
- *reduce the stride of internal checkpoints*: to guarantee that soft exits leave the simulation in a state that can be restarted safely, i-PI stores the internal state of the simulation at each step. Particularly for complicated setups, the overhead can be substantial. You can reduce the frequency by which the internal state is stored using the *safe_stride* attribute in the `<simulation>` tag. Note that doing so increases the risk that the RESTART file saved upon soft exit will be inconsistent with the state of the outputs, so that restarting a simulation will leave broken or discontinuous output files.
- *profile i-PI*: if you still think i-PI is being unreasonably slow, you can contact the developers - your setup might have revealed some kind of bottleneck. It will help us if you can also generate a profiler output from your run (possibly with only a small number of steps). You can generate a profiler log by running i-PI with *-p* option, e.g. *i-pi -p input.xml*. You will need to install the *yappi* profiler.

8 Frequently asked questions

8.1 Where to ask for help?

There is distinction between questions and bugs. If your question is related to running i-PI, using or understanding its features (also if you think that something is not documented enough), the best way to ask is to post your question on the i-PI forum: <https://groups.google.com/g/ipi-users>. In the forum, everything discussed will be available to other members of the community in the future!

If you are seeing an actual bug, please consider reporting an issue on Github: <https://github.com/i-pi/i-pi/issues>.

8.2 Which units does i-PI use?

Atomic units are used everywhere inside the code. The exact values are listed in the manual section “Internal units”. For I/O operations, various units are accepted. The XML tags responsible for physical settings accept the attribute `units=<unit>`. If nothing is specified, i-PI will assume atomic units. Many of the provided examples have these tags, and the reference section of the manual contains a full list. Another way to use units is related to input `.xyz` files. i-PI can accept declarations of these units in the `.xyz` comment line. One example is:

```
# CELL(abcABC): a b c alpha beta gamma cell{units} positions/forces/velocities{units}
```

Another place worth looking at is `<i-pi-root>/ipi/utls/units.py`, which contains definitions and values for non-default units, such as `electronvolt`, `inversecm` etc.

8.3 How to build i-PI?

i-PI is a Python code, and as such, strictly speaking, does not need to be compiled and installed. `<i-pi-root>/bin/i-pi` file is an executable. `source <i-pi-root>/env.sh` ensures that necessary paths are added to `PATH`, `PYTHONPATH` etc. However, it is often more convenient to install the package to the system’s Python modules path, so that it is accessible by all users and can be run without specifying the path to the Python script. For this purpose we have included a module in the root directory of the i-PI distribution, `setup.py`, which handles creating a package with the executable and all the modules which are necessary for it to run. Detailed explanation can be found in the manual.

In addition, i-PI can be installed through `pip` with the command

```
$ pip install -U i-PI
```

And to test it with `pytest`

```
$ pip install pytest
$ pytest --pyargs ipi.tests
```

8.4 How to run i-PI with the client code <name>?

i-PI communicates with electronic structure and MD codes via socket communication protocol. For many popular codes (CP2K, FHI-aims, LAMMPS, Quantum ESPRESSO, VASP, etc.) we provide examples in `<i-pi-root>/examples` folder.

Another way of connecting to client codes is using the ASE client. This way you get access to wide variety of codes that are connected to ASE, but for some of them current implementation requires restarting the client code after each step, which may lead to a significant overhead in case of electronic structure calculations. We recommend using the socket connection from the client code to ASE and then from ASE to i-PI, when possible. An example of a “double-socket” setup can be found in:

<https://github.com/i-pi/i-pi/tree/master/examples/ASEClient>

8.5 How to run i-PI on a cluster?

There are different ways of running i-PI on HPC systems, described in details in the manual. One simple setup, which assumes that i-PI and all client codes are launched from the same Slurm script, is given in `examples/slurm/sl-rpc`. In this example, i-PI is launched in the background first, and then its hostname is passed to the input files of the clients. Clients can communicate to the same socket (identified by the same port number) or to different sockets (different port numbers) that i-PI has opened. The clients are subsequently launched from the same slurm script.

Other setups are possible, please consult the manual.

8.6 How to setup colored-noise thermostats with meaningful parameters?

If you have an application that will profit from colored noise thermostats, you can download the parameters from <http://gle4md.org>. This website contains sets of input parameters optimized for a variety of conditions. You can choose i-PI syntax, such that you will only have to copy-paste that section inside the i-PI input file. Also the `<i-pi-root>/examples` (especially withing `lammps`) folder demonstrates the syntax for various thermostats.

8.7 How to perform geometry optimization?

Several examples of geometry optimizaion, including fixing certain atoms, are present in the `<i-pi-root>/examples` folder. The folders with names `*geop*` correspond to geometry optimization. Note that the optimization algorithms that involve line search have two sets of convergence parameters: one for line search and another for the main algorithm.

9 Contributing

i-PI is an open source project, and everyone is welcome to contribute with bug fixes, documentation, examples, new features, regression tests, etc. Any contribution from users or developers is very much welcome! These contributions may be full new methods/algorithms that you have developed or that you just would like to add to i-PI, or simple improvements (that would make your life easier) regarding the code itself, the documentation, or analysis scripts.

9.1 Useful resources

For more information on i-PI in general please consult the following places:

- i-PI homepage: <http://ipi-code.org>
- i-PI forum: <https://groups.google.com/g/ipi-users>
- i-PI github repo, including issues page: <https://github.com/i-pi/i-pi>

9.2 What we expect from a contribution

To start contributing, note the following general guidelines:

- Your contribution should be based on the master branch of the repository
- We expect you to fork the repository and make all the changes in your fork
- When done, you can open a pull request (PR) to our repository, which will be reviewed.

For a PR to be issued, we will expect that:

- You have produced code that is compliant with our formatting requirements (*black* and *flake8*, as explained in the README here: <https://github.com/i-pi/i-pi>)
- You have written understandable docstrings for all functions and classes you have created
- For all contributions (bug fixes, new functionalities, etc.), ensure that the continuous integration tests all pass.
- In case of a new functionality:
 - You have added an example to the *examples* folder that showcases it. The inputs should be a simple but meaningful example of the functionality.
 - You have added a regression test to *i-pi/ipi_tests/regression_tests* folder (more on this below).

9.3 Creating a regression test

If you code up a new functionality, it is best that it goes directly to the regression test infrastructure, so that we can ensure that it does not break with future code restructuring or the addition of new features in the future. The information below can also be found in the README file in the *i-pi/ipi_tests/regression_tests* folder:

To add a new regression test please provide:

- *input.xml* (and all the files required by it)
- *test_settings.dat* This file specifies the *driver_model* and flags that will be used when running the driver to perform the regression test. For examples including the usage of option flags, please see:
`tests/NVE/NVE_1/harmonic_python/test_settings.dat` `tests/NVE/NVE_1/harmonic/test_settings.dat`
- *file_to_check.txt* specifying the files serving as reference with their respective filenames. For an existing example, please see:
`tests/INSTANTON/100K/files_to_check.txt`
- reference files that are listed in the *files_to_check.txt* file. Currently, the available formats are *.xyz* and *numpy-accessible* which should be specified with the keywords ‘xyz’ and ‘numpy’.

Important points to note:

- In the case of multiple drivers needed for a simulation, the *test_settings.dat* file has to include a line for each driver by specifying as many lines as needed of *driver_model xxxx* in the order they should be assigned to the sockets. That is followed by lines describing the corresponding flags of each driver.
- The extension **.out* appears in the *.gitignore* file. This means that the ‘ref_simulation.out’ file has to be added manually to your commit. You can do that by typing: “git add -f ref_simulation.out”

If you are developing something new and you are unsure about how to proceed, we encourage opening an issue with your question. Depending on the development, we can open a channel of discussion with the core developers (Slack) in order to see how to best handle your case!

9.4 Adding features involving a client code

If your new development in i-PI is directly related to a specific client code or if there is the need of coding a new native client in i-PI there can be some further adjustments to do. We very much welcome new interfaces and we will be happy to answer your questions.

If you want to enable the communication of a new client code with i-PI, it is not difficult: Please check an example of how it was done in *fortran* and *python* in the *drivers* folder in the repository.

9.5 Getting recognition for your contribution

We recognize that adding new features or enhancements to any project is a task that needs more recognition. If you want your contribution advertised, we are happy to do so through our webpage. Just let us know and we will add your name, a short description, and links to any publications [here](#).

9.6 Contributing with a bug report

Bug reports are a very welcome feedback! These should be posted in the i-PI Github [repository](#) (issues tab). Please include inputs and outputs that allow us to reproduce the bug, and a short description of the bug you are encountering.

9.7 Questions related to i-PI usage

Questions related to usage or general questions are best posted in our forum (see link above to google groups). If we establish that there is the need to open a Github issue related to the question, we will do so.

10 Bibliography

References

- [BM06] Bastiaan J Braams and David E Manolopoulos. On the short-time limit of ring polymer molecular dynamics. *J. Chem. Phys.*, 125:124105, 2006.
- [BDP07] G Bussi, D Donadio, and M Parrinello. Canonical sampling through velocity rescaling. *J. Chem. Phys.*, 126:14101, 2007.
- [CV93] Jianshu Cao and Gregory A Voth. A new perspective on quantum time correlation functions. *J. Chem. Phys.*, 99:10070–10073, 1993.
- [CV94] Jianshu Cao and Gregory A Voth. The formulation of quantum statistical mechanics based on the Feynman path centroid density. IV. Algorithms for centroid molecular dynamics. *J. Chem. Phys.*, 101:6168–6183, 1994.
- [Cer10] M. Ceriotti. GLE4MD. <http://gle4md.org>, 2010.
- [CBRM11] Michele Ceriotti, Guy A. R. Brain, Oliver Riordan, and David E. Manolopoulos. The inefficiency of re-weighted sampling and the curse of system size in high order path integration. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 468:2–17, 2011.
- [CBP09a] Michele Ceriotti, Giovanni Bussi, and Michele Parrinello. Langevin Equation with Colored Noise for Constant-Temperature Molecular Dynamics Simulations. *Phys. Rev. Lett.*, 102:020601, 2009.
- [CBP09b] Michele Ceriotti, Giovanni Bussi, and Michele Parrinello. Nuclear Quantum Effects in Solids Using a Colored-Noise Thermostat. *Phys. Rev. Lett.*, 103:30603, 2009.
- [CBP10] Michele Ceriotti, Giovanni Bussi, and Michele Parrinello. Colored-Noise Thermostats à la Carte. *J. Chem. Theory Comput.*, 6:1170–1180, 2010.
- [CM12] Michele Ceriotti and David E Manolopoulos. Efficient First-Principles Calculation of the Quantum Kinetic Energy and Momentum Distribution of Nuclei. *Phys. Rev. Lett.*, 109:100604, 2012.

- [CMP11] Michele Ceriotti, David E Manolopoulos, and Michele Parrinello. Accelerating the convergence of path integral dynamics with a generalized Langevin equation. *J. Chem. Phys.*, 134:84104, 2011.
- [CM13] Michele Ceriotti and Thomas E Markland. Efficient methods and practical guidelines for simulating isotope effects. *J. Chem. Phys.*, 138:014112, 2013.
- [CMM14] Michele Ceriotti, Joshua More, and David E. Manolopoulos. i-PI: A Python interface for ab initio path integral molecular dynamics simulations. *Comp. Phys. Comm.*, 185:1019–1026, 2014.
- [CP10] Michele Ceriotti and Michele Parrinello. The δ -thermostat: selective normal-modes excitation by colored-noise Langevin dynamics. *Procedia Computer Science*, 1:1607–1614, 2010.
- [CPMM10] Michele Ceriotti, Michele Parrinello, Thomas E Markland, and David E Manolopoulos. Efficient stochastic thermostating of path integral molecular dynamics. *J. Chem. Phys.*, 133:124104, 2010.
- [CBC16] Bingqing Cheng, Jörg Behler, and Michele Ceriotti. Nuclear Quantum Effects in Water at the Triple Point: Using Theory as a Link Between Experiments. *J. Phys. Chem. Letters*, 7:2210–2215, 2016.
- [CC14] Bingqing Cheng and Michele Ceriotti. Direct path integral estimators for isotope fractionation ratios. *J. Chem. Phys.*, 141:244112, 2014.
- [Chi97] Siu A. Chin. Symplectic integrators from composite operator factorizations. *Phys. Lett. A*, 226:344–348, 1997.
- [CM04] I R Craig and D E Manolopoulos. Quantum statistics and classical mechanics: Real time correlation functions from ring polymer molecular dynamics. *J. Chem. Phys.*, 121:3368, 2004.
- [FH64] R P Feynman and A R Hibbs. *Quantum Mechanics and Path Integrals*. McGraw-Hill, New York, 1964.
- [HMMM13] Scott Habershon, David E Manolopoulos, Thomas E Markland, and Thomas F Miller. Ring-polymer molecular dynamics: quantum effects in chemical dynamics from classical trajectories in an extended phase space. *Annual review of physical chemistry*, 64:387–413, 2013.
- [HWC18] Mahdi Hijazi, David M. D.M. Wilkins, and Michele Ceriotti. Fast-forward Langevin dynamics with momentum flips. *J. Chem. Phys.*, 148:184109, 2018.
- [JV01] Seogjoo Soonmin Jang and Gregory A Voth. Applications of higher order composite factorization schemes in imaginary time path integral simulations. *J. Chem. Phys.*, 115:7832–7842, 2001.
- [KBC16] Venkat Kapil, Jörg Behler, and Michele Ceriotti. High order path integrals made easy. *J. Chem. Phys.*, 145:234103, 2016.
- [KCC18] Venkat Kapil, Alice Cuzzocrea, and Michele Ceriotti. Anisotropy of the Proton Momentum Distribution in Water. *J. Phys. Chem. B*, 122:6048–6054, 2018.
- [KRM+19] Venkat Kapil, Mariana Rossi, Ondrej Marsalek, Riccardo Petraglia, Yair Litman, Thomas Spura, Bingqing Cheng, Alice Cuzzocrea, Robert H Meißner, David M Wilkins, Benjamin A Helfrecht, Przemysław Juda, Sébastien P Bienvenue, Wei Fang, Jan Kessler, Igor Poltavsky, Steven Vandenbrande, Jelle Wieme, Clemence Corminboeuf, Thomas D Kühne, David E Manolopoulos, Thomas E Markland, Jeremy O Richardson, Alexandre Tkatchenko, Gareth A Tribello, Veronique Van Speybroeck, and Michele Ceriotti. i-PI 2.0: A universal force engine for advanced molecular simulations. *Computer Physics Communications*, 236:214–223, March 2019.
- [KVC16] Venkat Kapil, Joost VandeVondele, and Michele Ceriotti. Accurate molecular dynamics and nuclear quantum effects at low cost by multiple steps in real and imaginary time: Using density functional theory to accelerate wavefunction methods. *J. Chem. Phys.*, 144:054111, 2016.
- [LMCP10] Lin Lin, Joseph A Morrone, Roberto Car, and Michele Parrinello. Displaced Path Integral Formulation for the Momentum Distribution of Quantum Particles. *Phys. Rev. Lett.*, 105:110602, 2010.
- [LDCR18] Yair Litman, Davide Donadio, Michele Ceriotti, and Mariana Rossi. Decisive role of nuclear quantum effects on surface mediated water dissociation at finite temperature. *J. Chem. Phys.*, 148:102320, 2018.

- [LAM+13] Jian Liu, Richard S. Andino, Christina M. Miller, Xin Chen, David M. Wilkins, Michele Ceriotti, and David E. Manolopoulos. A Surface-Specific Isotope Effect in Mixtures of Light and Heavy Water. *J. Phys. Chem. C*, 117:2944–2951, 2013.
- [MM08a] Thomas E Markland and David E Manolopoulos. A refined ring polymer contraction scheme for systems with electrostatic interactions. *Chem. Phys. Lett.*, 464:256, 2008.
- [MM08b] Thomas E Markland and David E Manolopoulos. An efficient ring polymer contraction scheme for imaginary time path integral simulations. *J. Chem. Phys.*, 129:024105, 2008.
- [MHT99] Glenn J Martyna, Adam Hughes, and Mark E Tuckerman. Molecular dynamics algorithms for path integrals at constant pressure. *J. Chem. Phys.*, 110:3275, 1999.
- [MM05] Thomas F Miller and David E Manolopoulos. Quantum diffusion in liquid para-hydrogen from ring-polymer molecular dynamics. *J. Chem. Phys.*, 122:184503, 2005.
- [PNicolaiW+16] Riccardo Petraglia, Adrien Nicolaï, M.D. Matthew D. Wodrich, Michele Ceriotti, and C. Corminboeuf. Beyond static structures: Putting forth REMD as a tool to solve problems in computational organic chemistry. *J. Comp. Chem.*, 37:83–92, 2016.
- [PT16] Igor Poltavsky and Alexandre Tkatchenko. Modeling quantum nuclei with perturbed path integral molecular dynamics. *Chemical Science*, 7:1368–1372, 2016.
- [PerezT11] Alejandro Pérez and Mark E Tuckerman. Improving the convergence of closed and open path integral molecular dynamics via higher order Trotter factorization schemes. *J. Chem. Phys.*, 135:064104, 2011.
- [RCM14] Mariana Rossi, Michele Ceriotti, and David E Manolopoulos. How to remove the spurious resonances from ring polymer molecular dynamics. *J. Chem. Phys.*, 140:234116, 2014.
- [RGC16] Mariana Rossi, Piero Gasparotto, and Michele Ceriotti. Anharmonic and Quantum Fluctuations in Molecular Crystals: A First-Principles Study of the Stability of Paracetamol. *Phys. Rev. Lett.*, 117:115702, 2016.
- [RKC18] Mariana Rossi, Venkat Kapil, and Michele Ceriotti. Fine tuning classical and quantum molecular dynamics using a generalized Langevin equation. *J. Chem. Phys.*, 148:102301, 2018.
- [SG78] Isaac F. Silvera and Victor V. Goldman. The isotropic intermolecular potential for H₂ and D₂ in the solid and gas phases. *J. Chem. Phys.*, 69:4209, 1978.
- [Suz95] M Suzuki. Hybrid exponential product formulas for unbounded operators with possible applications to Monte Carlo simulations. *Phys. Lett. A*, 201:425–428, 1995.
- [Yam05] Takeshi M Yamamoto. Path-integral virial estimator based on the scaling of fluctuation coordinates: Application to quantum clusters with fourth-order propagators. *J. Chem. Phys.*, 123:104101, 2005.